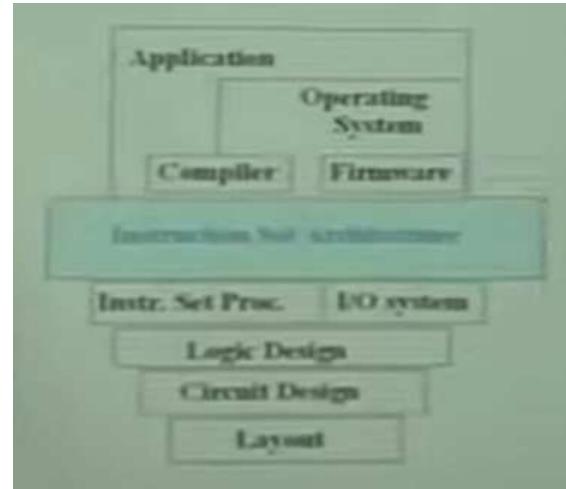
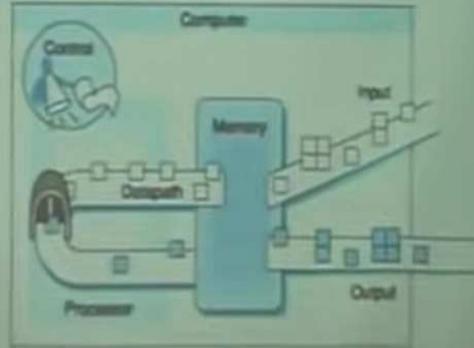


Introduction to Computer Organization _ Basic Concept

Computer Organization and Architecture

- Five classic components of a computer – input, output, memory, datapath, and control



The Instruction Set: a Critical Interface



ISA (Instruction Set Architecture)

- ISA, or simply architecture – the abstract interface between the hardware and the lowest level software that encompasses all the information necessary to write a machine language program, including instructions, registers, memory access, I/O, ...
 - Enables implementations of varying cost and performance to run identical software
- The combination of the basic instruction set (the ISA) and the operating system interface is called the application binary interface (ABI)
 - ABI – The user portion of the instruction set plus the operating system interfaces used by application programmers. Defines a standard for binary portability across computers.

- "... the attributes of a [computing] system as seen by the programmer, i.e. the conceptual structure and functional behavior, as distinct from the organization of the data flow and controls, the logic design, and the physical implementation."
- Amdahl, Blaauw, and Brooks, 1964
- ISA Includes
 - Organization of storage
 - Data types
 - Encoding and representing instructions
 - Instruction Set (or opcodes)
 - Modes of addressing data items/instructions
 - Program visible exception handling
- Specifies requirements for binary compatibility across implementations (ABI)

- A very important abstraction
 - interface between hardware and low-level software
 - standardizes instructions, machine language bit patterns, etc
 - advantage: *different implementations of the same architecture*
 - disadvantage: *sometimes prevents using new innovations*
- Common instruction set architectures:
 - IA-32, PowerPC, MIPS, SPARC, ARM, and others

Introduction to Computer Organization _ MIPS Instruction and Datapath

MIPS Instruction and Format

MIPS instructions	Name	Format	Pseudo MIPS	Name	Format
add	add	R	move	move	R
subtract	sub	R	multiply	mult	R
add immediate	addi	I	multiply immediate	multi	I
load word	lw	I	load immediate	li	I
store word	sw	I	branch less than	blt	I
load half	lh	I	branch less than		
load half unsigned	lhu	I	or equal	ble	I
store half	sh	I	branch greater than	bgt	I
load byte	lb	I	branch greater than		
load byte unsigned	lbu	I	or equal	bge	I
store byte	sb	I			
load linked	ll	I			
store conditional	sc	I			
load upper immediate	lui	I			
and	and	R			
or	or	R			
nor	nor	R			
and immediate	andi	I			
or immediate	ori	I			
shift left logical	sll	R			
shift right logical	srl	R			
branch on equal	beq	I			
branch on not equal	bne	I			
set less than	slt	R			
set less than immediate	slti	I			
set less than immediate unsigned	sltiu	I			
jump	j	J			
jump register	jr	R			
jump and link	jal	J			

Instruction and Datapath

	Instruction name	ARM	MIPS
Register-register	Add	add	addu, addiu
	Add (trap if overflow)	adds; swivs	add
	Subtract	sub	subu
	Subtract (trap if overflow)	subs; swivs	sub
	Multiply	mul	mult, multu
	Divide	—	div, divu
	And	and	and
	Or	orr	or
	Xor	eor	xor
	Load high part register	—	lui
	Shift left logical	lsl ¹	sllv, sll
	Shift right logical	lsr ¹	srlv, srl
Data transfer	Shift right arithmetic	asr ¹	sra, sra
	Compare	cmp, cmn, tst, teq	slt/i,slt/iu
	Load byte signed	ldrsb	lb
	Load byte unsigned	ldrb	lbu
	Load halfword signed	ldrsh	lh
	Load halfword unsigned	ldrh	lhu
	Load word	ldr	lw
	Store byte	strb	sb
	Store halfword	strh	sh
	Store word	str	sw
	Read, write special registers	mrs, msr	move
	Atomic Exchange	swp, swpb	ll;sc

FIGURE 2.32 ARM register-register and data transfer instructions equivalent to MIPS core. Dashes mean the operation is not available in that architecture or not synthesized in a few instructions. If there are several choices of instructions equivalent to the MIPS core, they are separated by commas. ARM includes shifts as part of every data operation instruction, so the shifts with superscript 1 are just a variation of a move instruction, such as lsl^1 . Note that ARM has no divide instruction.

Introduction to Computer Organization _ MIPS (RISC) Design Principle

1. Simplicity favors regularity

Arithmetic Operations

- Add and subtract, three operands

- | Two sources and one destination

```
add a, b, c # a gets b + c
```

- All arithmetic operations have this form

- *Design Principle 1: Simplicity favors regularity*

- | Regularity makes implementation simpler
 - | Simplicity enables higher performance at lower cost

2. Smaller is faster

Register Operands

- Arithmetic instructions use register operands

- MIPS has a 32×32 -bit register file
 - | Use for frequently accessed data
 - | Numbered 0 to 31
 - | 32-bit data called a "word"

- Assembler names

- | \$t0, \$t1, ..., \$t9 for temporary values
 - | \$s0, \$s1, ..., \$s7 for saved variables

- *Design Principle 2: Smaller is faster*

- | c.f. main memory: millions of locations

4. Good Design demands good compromises

Branch Instruction Design

- Why not include `blt`, `bge`, etc in the MIPS ISA?

- Hardware for $<$, \geq , ... slower than $=$, \neq

- | Combining with branch involves more work per instruction, requiring a slower clock
 - | All instructions would be penalized!

- `beq` and `bne` are the common case (Principle #3)

- This is a good design compromise (Principle #4)

3. Make the common case fast

Immediate Operands

- Constant data specified in an instruction

```
addi $s3, $s3, 4
```

- No subtract immediate instruction

- | Just use a negative constant

```
addi $s2, $s1, -1
```

- *Design Principle 3: Make the common case fast*

- | Small constants are common (50% of MIPS arithmetic instructions in SPEC2006 use constants!)
 - | Immediate operand avoids a load instruction

Introduction to Computer Organization _ MIPS Addressing Mode

MIPS Instruction Format

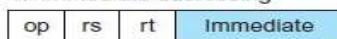
Name	Bit Fields						Notes
	6 bits	5 bits	5 bits	5 bits	6 bits	(32 bits total)	
R-Format	op	rs	rt	rd	shmt	funct	Arithmetic, logic
I-format	op	rs	rt	address/immediate (16)			Load/store, branch, immediate
J-format	op	target address (26)					Jump

Popularity

Instruction class	MIPS examples	HLL correspondence	Frequency	
			Integer	Ft. pt.
Arithmetic	add, sub, addi	Operations in assignment statement s	16%	48%
Data transfer	lw, sw, lb, lh, lhu, sb, lui	References to data structures, such as arrays	35%	36%
Logical	and, or, nor, andi, ori, sll, srl	Operations in assignment statement s	12%	4%
Conditional branch	beq, bne,slt, slti, sltiu	If statements and loops	34%	8%
Jump	j, jr, jal	Procedure calls, returns, and case/switch statements	2%	0%

MIPS Addressing Mode

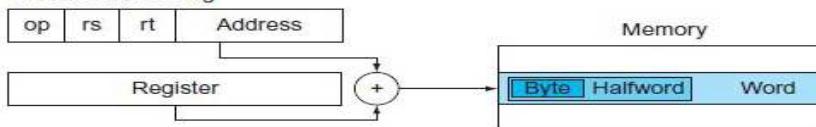
1. Immediate addressing



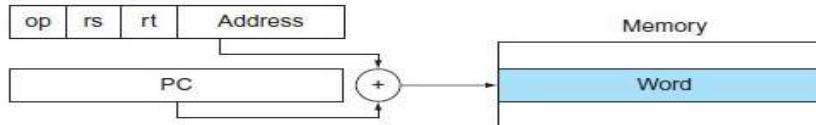
2. Register addressing



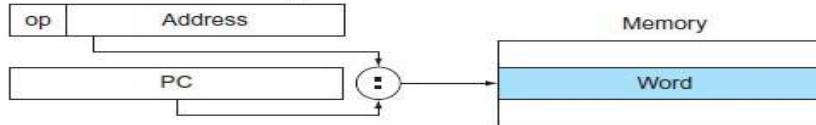
3. Base addressing



4. PC-relative addressing



5. Pseudodirect addressing



MIPS Instruction Category

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	add \$s1,\$s2,\$s3	\$s1 = \$s2 + \$s3	Three register operands
	subtract	sub \$s1,\$s2,\$s3	\$s1 = \$s2 - \$s3	Three register operands
	add immediate	addi \$s1,\$s2,20	\$s1 = \$s2 + 20	Used to add constants
	load word	lw \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Word from memory to register
	store word	sw \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Word from register to memory
	load half	lh \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Halfword memory to register
	load half unsigned	lhu \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Halfword memory to register
	store half	sh \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Halfword register to memory
	load byte	lb \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Byte from memory to register
	load byte unsigned	lbu \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Byte from memory to register
Data transfer	store byte	sb \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Byte from register to memory
	load linked word	l1 \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Load word as 1st half of atomic swap
	store condition. word	sc \$s1,20(\$s2)	Memory[\$s2+20]=1:\$s1=0 or 1	Store word as 2nd half of atomic swap
	load upper immmed.	lui \$s1,20	\$s1 = 20 * 2 ¹⁶	Loads constant in upper 16 bits
	and	and \$s1,\$s2,\$s3	\$s1 = \$s2 & \$s3	Three reg. operands; bit-by-bit AND
	or	or \$s1,\$s2,\$s3	\$s1 = \$s2 \$s3	Three reg. operands; bit-by-bit OR
	nor	nor \$s1,\$s2,\$s3	\$s1 = ~(\$s2 & \$s3)	Three reg. operands; bit-by-bit NOR
	and immediate	andi \$s1,\$s2,20	\$s1 = \$s2 & 20	Bit-by-bit AND reg with constant
	or immediate	ori \$s1,\$s2,20	\$s1 = \$s2 20	Bit-by-bit OR reg with constant
	shift left logical	sll \$s1,\$s2,10	\$s1 = \$s2 << 10	Shift left by constant
Logical	shift right logical	srl \$s1,\$s2,10	\$s1 = \$s2 >> 10	Shift right by constant
	branch on equal	beq \$s1,\$s2,25	if(\$s1 == \$s2) go to PC + 4 + 100	Equal test; PC-relative branch
	branch on not equal	bne \$s1,\$s2,25	if(\$s1 != \$s2) go to PC + 4 + 100	Not equal test; PC-relative
	set on less than	slt \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than; for beq, bne
	set on less than unsigned	sltu \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than unsigned
	set less than immediate	slti \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant
	set less than immediate unsigned	sltiu \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant unsigned
Conditional branch	set on less than	slt \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than; for beq, bne
	set on less than immediate	slti \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant
	set less than immediate	sltiu \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant unsigned
	set less than immediate unsigned	sltiu \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant unsigned

Introduction to Computer Organization _ MIPS Program Execution

MIPS Instruction Executions

- Data operations**
 - Arithmetic (add, subtract, ...)
 - Logical (and, or, not, xor, ...)
- Data transfer**
 - Load (memory \rightarrow register)
 - Store (register \rightarrow memory)
- Sequencing**
 - Branch (conditional, e.g., $<$, $>$, $==$)
 - Jump (unconditional, e.g., goto)

Function	Instruction	Effect
add	add R1, R2, R3	R1 = R2 + R3
sub	sub R1, R2, R3	R1 = R2 - R3
add immediate	addi R1, R2, 145	R1 = R2 + 145
multiply	mult R1, R2, R3	remainder
divide	div R1, R2, R3	R1 = R2 / R3
and	and R1, R2, R3	R1 = R2 & R3
or	or R1, R2, R3	R1 = R2 R3
and immediate	andi R1, R2, 145	R1 = R2 & 145
or immediate	ori R1, R2, 145	R1 = R2 145
shift left logical	sll R1, R2, 7	R1 = R2 << 7
shift right logical	srl R1, R2, 7	R1 = R2 >> 7
load word	lw R1, 145(R2)	R1 = [R2 + 145]
store word	sw R1, 145(R2)	[R2 + 145] = R1
load upper immediate	lui R1, 145	R1 = 145
branch on equal	beq R1, R2, 145	If (R1 == R2) go to PC + 4 + 145*4
branch on not equal	bne R1, R2, 145	If (R1 != R2) go to PC + 4 + 145*4
set on less than	slt R1, R2, R3	If (R2 < R3) R1 = 1, else R1 = 0
set less than immediate	slti R1, R2, 1	R1 = 1, else R1 = 0
jump	jr R1, 145	R31 = PC + 4; go to 145
jump register	j R1	R31 = PC + 4; go to R1
jump and link	jal 145	R31 = PC + 4; go to 145

(Complete table is printed in the book for reference.)

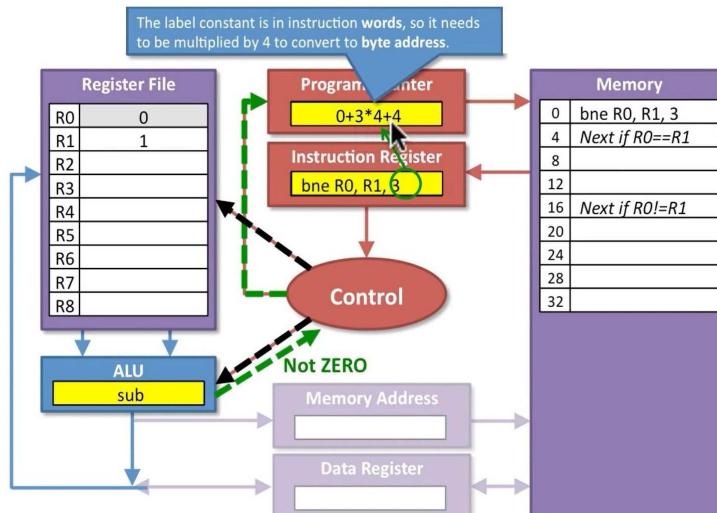
1. Data operations:
add/sub

2. Data transfers:
load word/store word

3. Sequencing:
branch equal/jump

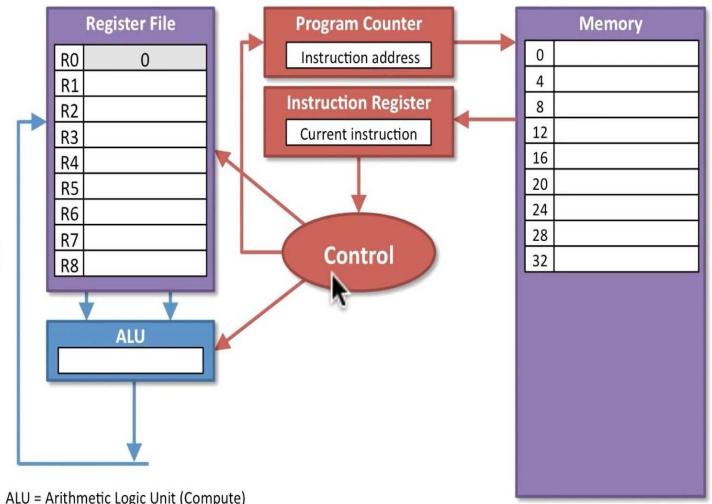
3. Sequencing

- ALU compares registers
- Result tells the Control whether to branch
- If the branch is taken, then the Control adds a constant from the instruction to the Program Counter
- The Control always adds 4 to the Program Counter



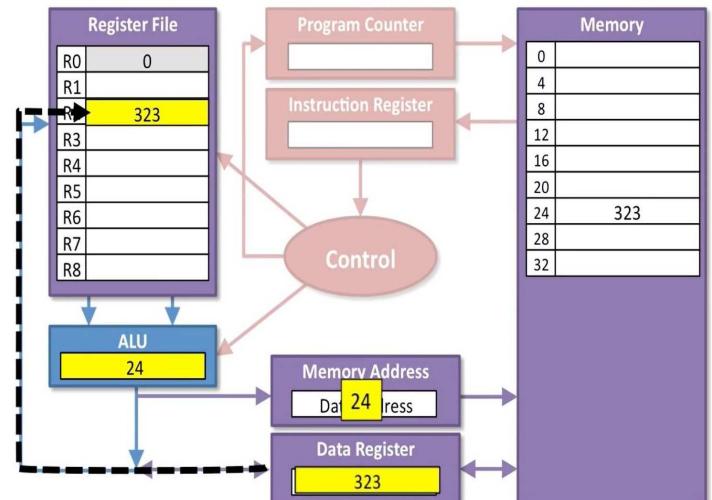
1. Data Operations

- Program Counter holds the instruction address.
- Instructions are fetched from memory into the Instruction Register.
- Control logic decodes the instruction and tells the ALU and Register File what to do.
- ALU executes the instruction and results flow back to the Register File.
- The Control logic updates the Program Counter for the next instruction.



2. Data Transfer

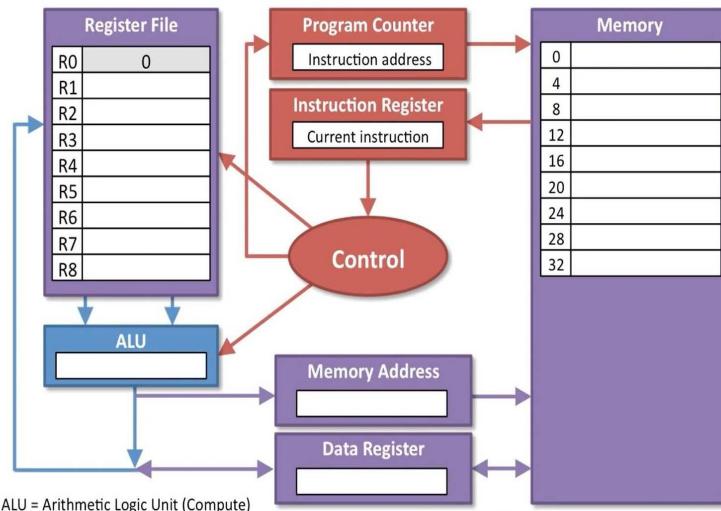
- ALU generates address
- Address goes to the Memory Address Register
- Results to/from memory are stored in the Memory Data Register
- Data from memory can now be stored back into the Register File or to memory can be written.



Introduction to Computer Organization _ MIPS Branch and Jump

Execution Review

1. Program Counter holds the instruction address.
2. Instructions are **fetched** from memory into the Instruction Register.
3. Control logic **decodes** the instruction and tells the ALU and Register File what to do.
4. ALU **executes** the instruction and results flow back to the Register File.
5. The Control logic **updates** the Program Counter for the next instruction.
6. The Memory Address register and Data Register are used to load and store to/from Memory.



Constants(immediate)

- Small **constants (immediates)** are used all over code (~50%)

```
if (a==b) c=1;
else c=2;
```

- How can we support this in the processor?

- Put the “typical constants” in memory and load them (slow)
- Create hard-wired registers (like R0) for them (how many?)

Need lots of bits to choose among lots of constant registers.

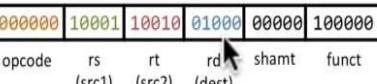
- MIPS does something in between:

- Some instructions can have **constants inside the instruction**
- The control logic then sends the constants to the ALU
- **addi R29, R29, 4** ← value 4 is inside the instruction

Store the constant data in the instruction, not the register file.

- But there's a problem:

- Instructions have only 32 bits. Need some for opcode and registers.



- How do we tradeoff space for constants and instructions?

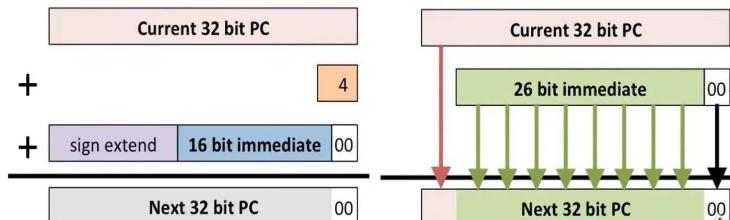
Branch and Jump

Branch instructions

- bne/beq I-format **16 bit immediate**
- j J-format **26 bit immediate**

But addresses are 32 bits! How do we handle this?

- Treat bne/beq as **relative offsets (add to current PC)**
- Treat j as an **absolute value (replace 26 bits of the PC)**

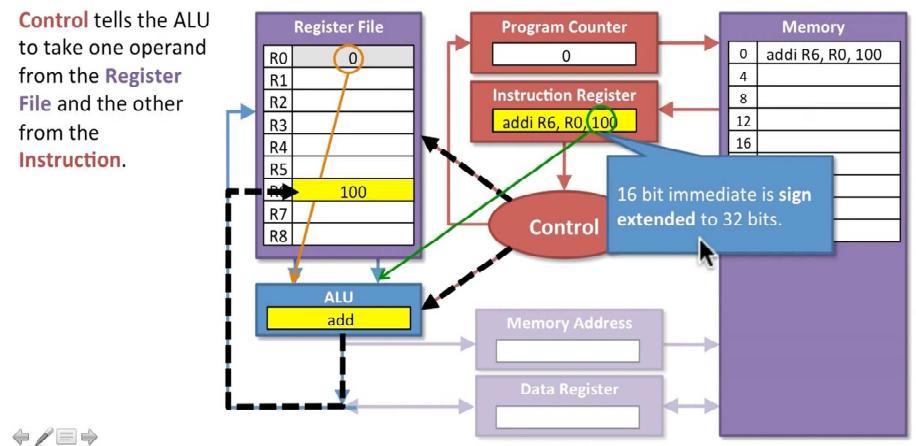


Loading immediate

Loading immediate values (constants)

14

Control tells the ALU to take one operand from the Register File and the other from the Instruction.



Introduction to Computer Organization _ MIPS ISA (Instruction Set Architecture)

MIPS instructions	Name	Format	Pseudo MIPS	Name	Format
add	add	R	move	move	R
subtract	sub	R	multiply	mult	R
add immediate	addi	I	multiply immediate	multi	I
load word	lw	I	load immediate	li	I
store word	sw	I	branch less than	blt	I
load half	lh	I	branch less than or equal	ble	I
load half unsigned	lhu	I	branch greater than	bgt	I
store half	sh	I	branch greater than or equal	bge	I
load byte	lb	I			
load byte unsigned	lbu	I			
store byte	sb	I			
load linked	ll	I			
store conditional	sc	I			
load upper immediate	lui	I			
and	and	R			
or	or	R			
nor	nor	R			
and immediate	andi	I			
or immediate	ori	I			
shift left logical	sll	R			
shift right logical	srl	R			
branch on equal	beq	I			
branch on not equal	bne	I			
set less than	slt	R			
set less than immediate	slti	I			
set less than immediate	sltiu	I			
unsigned					
jump	j	J			
jump register	jr	R			
jump and link	jal	J			

MIPS Instruction Format							
Name	Bit Fields						Notes
	6 bits	5 bits	5 bits	5 bits	6 bits	(32 bits total)	
R-Format	op	rs	rt	rd	shmt	funct	Arithmetic, logic
I-format	op	rs	rt	address/immediate (16)			Load/store, branch, immediate
J-format	op	target address (26)					Jump

Logic Operation			
Logical operations	C operators	Java operators	MIPS instructions
Shift left	<<	<<	sll
Shift right	>>	>>>	srl
Bit-by-bit AND	&	&	and, andi
Bit-by-bit OR			or, ori
Bit-by-bit NOT	~	~	nor

FIGURE 2.8 C and Java logical operators and their corresponding MIPS instructions. MIPS implements NOT using a NOR with one operand being zero.

Basic Operations

- NOT
- AND
- OR
- NOR
- NAND
- XOR
- XNOR

Basic gates we will use in this class.

AND			OR			XOR		
A	B	Out	A	B	Out	A	B	Out
0	0	0	0	0	0	0	0	0
0	1	0	0	1	1	0	1	1
1	0	0	1	0	1	1	0	1
1	1	1	1	1	1	1	1	0

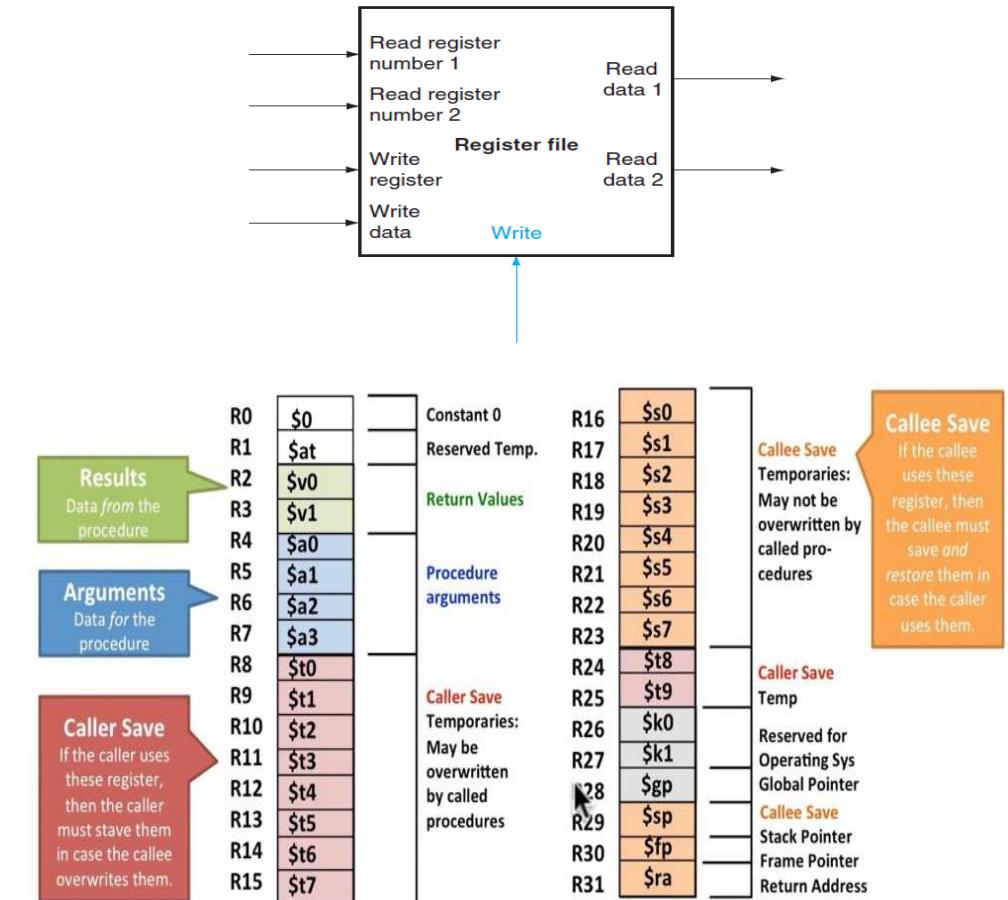
NAND			NOR			XNOR		
A	B	Out	A	B	Out	A	B	Out
0	0	1	0	0	1	0	0	1
0	1	0	0	1	0	0	1	0
1	0	1	1	0	0	1	0	0
1	1	0	1	1	0	1	1	1

Introduction to Computer Organization _ MIPS Assembly and Registers

MIPS Assembly Category

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	add \$s1,\$s2,\$s3	\$s1 = \$s2 + \$s3	Three register operands
	subtract	sub \$s1,\$s2,\$s3	\$s1 = \$s2 - \$s3	Three register operands
	add immediate	addi \$s1,\$s2,20	\$s1 = \$s2 + 20	Used to add constants
Data transfer	load word	lw \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Word from memory to register
	store word	sw \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Word from register to memory
	load half	lh \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Halfword memory to register
	load half unsigned	lhu \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Halfword memory to register
	store half	sh \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Halfword register to memory
	load byte	lb \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Byte from memory to register
	load byte unsigned	lbu \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Byte from memory to register
	store byte	sb \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Byte from register to memory
	load linked word	ll \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Load word as 1st half of atomic swap
	store condition. word	sc \$s1,20(\$s2)	Memory[\$s2+20]= \$s1; \$s1=0 or 1	Store word as 2nd half of atomic swap
Logical	load upper immmed.	lui \$s1,20	\$s1 = 20 * 2 ¹⁶	Loads constant in upper 16 bits
	and	and \$s1,\$s2,\$s3	\$s1 = \$s2 & \$s3	Three reg. operands; bit-by-bit AND
	or	or \$s1,\$s2,\$s3	\$s1 = \$s2 \$s3	Three reg. operands; bit-by-bit OR
	nor	nor \$s1,\$s2,\$s3	\$s1 = ~(\$s2 \$s3)	Three reg. operands; bit-by-bit NOR
	and immediate	andi \$s1,\$s2,20	\$s1 = \$s2 & 20	Bit-by-bit AND reg with constant
	or immediate	ori \$s1,\$s2,20	\$s1 = \$s2 20	Bit-by-bit OR reg with constant
	shift left logical	sll \$s1,\$s2,10	\$s1 = \$s2 << 10	Shift left by constant
	shift right logical	srl \$s1,\$s2,10	\$s1 = \$s2 >> 10	Shift right by constant
Conditional branch	branch on equal	beq \$s1,\$s2,25	if(\$s1 == \$s2) go to PC + 4 + 100	Equal test; PC-relative branch
	branch on not equal	bne \$s1,\$s2,25	if(\$s1 != \$s2) go to PC + 4 + 100	Not equal test; PC-relative
	set on less than	slt \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than; for beq, bne
	set on less than unsigned	sltu \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than unsigned
	set less than immediate	slti \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant
	set less than immediate unsigned	sltiu \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant unsigned

MIPS Registers and Usage Convention



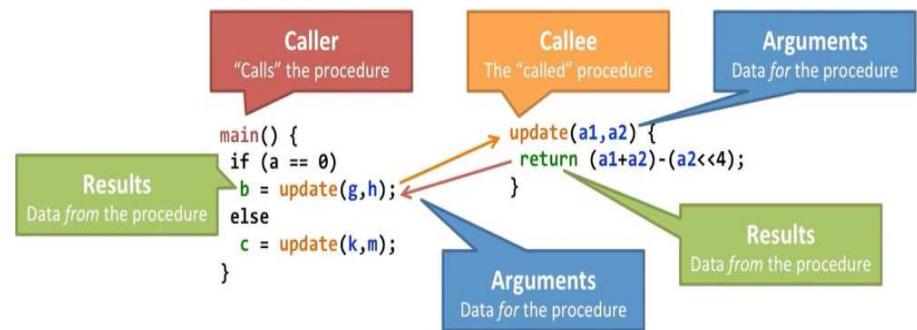
Name	Example	Comments
32 registers	\$s0-\$s7, \$t0-\$t9, \$zero, \$a0-\$a3, \$v0-\$v1, \$gp, \$fp, \$sp, \$ra, \$at	Fast locations for data. In MIPS, data must be in registers to perform arithmetic. Register \$zero always equals 0, and register \$at is reserved by the assembler to handle large constants.
2 ³⁰ memory words	Memory[0], Memory[4], ..., Memory[4294967292]	Accessed only by data transfer instructions. MIPS uses byte addresses, so sequential word addresses differ by 4. Memory holds data structures, arrays, and spilled registers.

Introduction to Computer Organization _ MIPS Procedure Calls 1

Procedure Calls

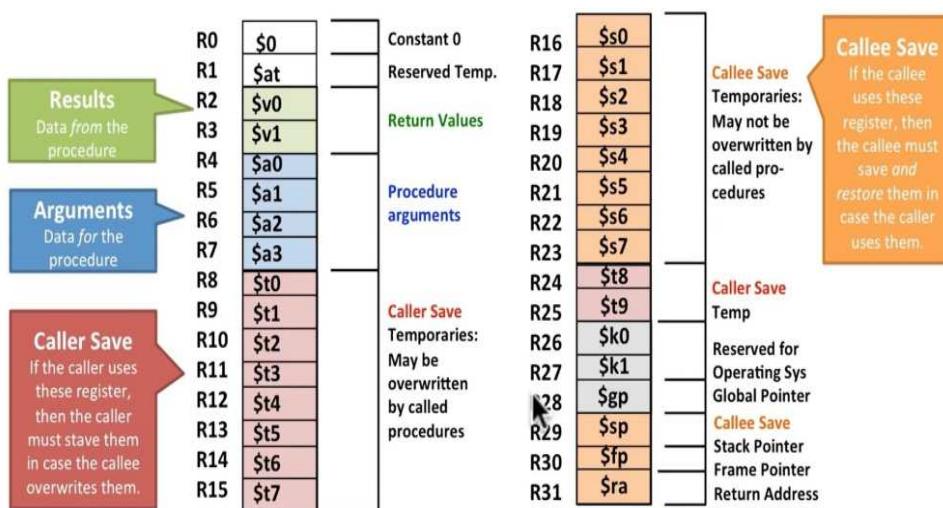
- Procedures need to:
 - Know where to find their **arguments** and put their **results**
 - Save and restore** registers to avoid overwriting the **caller's** registers
 - Return to the right place when done
- To accomplish this we:
 - Have **conventions** for who saves registers
 - Save them in memory on the **stack**
 - Use **jal** and **jr \$ra** to enter and exit procedures
- As long as everyone follows the convention we get interoperability

Terminology



- The **Caller** calls the procedure
- The **procedure** is the **Callee**
- The **Caller** gives the **Callee Arguments** (data)
- The **Callee** returns **Results** (data) to the **Caller**

Saving and Restoring Registers



How to do Procedure Call

- Transfer** control to the **callee** to start the procedure:

jal ProcedureAddress; jump-and-link to the procedure

- Keeps track of the instruction *after* the **jal** so we can continue in the right place when we are done with the procedure.
- Stores the return address (PC+4) in \$ra (R31)**

- Return** control to the **caller** when the procedure is done:

jr \$ra; jump-return to the address in \$ra

- Jumps back to the address stored in \$ra (R31)
- This is why you need to store the return address so you know where to go back to!

Introduction to Computer Organization _ MIPS Procedure Calls 1

6 Steps of Procedure Call

Six Steps in Execution of a Procedure

1. Main routine (caller) places parameters in a place where the procedure (callee) can access them
 - \$a0 - \$a3: four argument registers
2. Caller transfers control to the callee (`jal Dest`)
3. Callee acquires the storage resources needed
4. Callee performs the desired task
5. Callee places the result value in a place where the caller can access it
 - \$v0 - \$v1: two value registers for result values
6. Callee returns control to the caller (`jr $ra`)
 - \$ra: one return address register to return to the point of origin

Register Usage of Procedure Call

- \$a0 - \$a3: arguments (reg's 4 - 7)
- \$v0, \$v1: result values (reg's 2 and 3)
- \$t0 - \$t9: temporaries
 - Can be overwritten by callee
- \$s0 - \$s7: saved
 - Must be saved/restored by callee !
- \$gp: global pointer for static data (reg 28)
- \$sp: stack pointer (reg 29)
- \$fp: frame pointer (reg 30)
- \$ra: return address (reg 31)
- \$gp, \$sp, \$fp, \$ra must be saved/restored by callee !

Using Registers to Implement

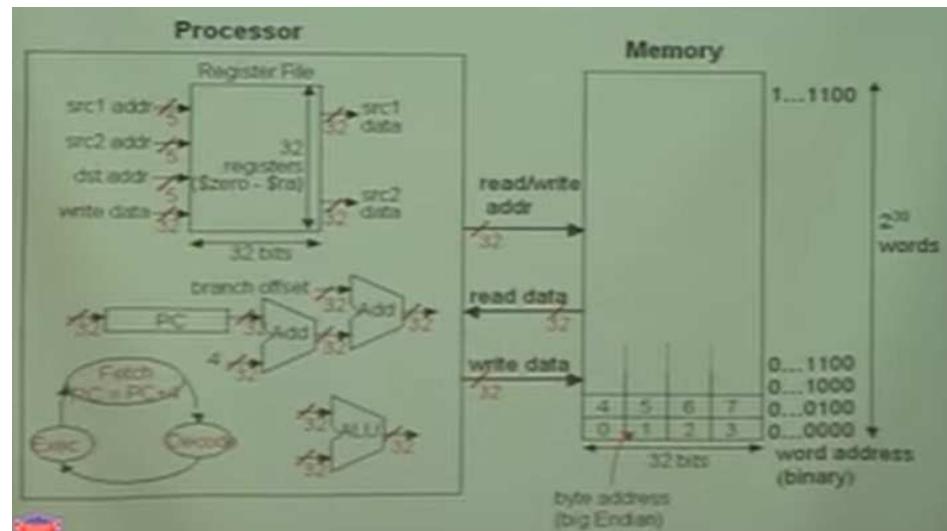
- Caller
 - Save caller-saved registers \$t0-\$t3, \$s0-\$s7
 - Load arguments in \$a0-\$a3, rest on stack above \$fp
 - Execute `jal` instruction
- Callee Setup
 - Allocate memory in frame ($$sp = ssp - \text{frame}$)
 - Save callee-saved registers \$s0-\$s7, \$fp, \$ra
 - Create frame ($$fp = ssp + \text{frame size} - 4$)
- Callee Return
 - Place return value in \$v0 and \$v1
 - Restore any callee-saved registers (\$fp, \$ra, \$s0-\$s7...)
 - Pop stack ($ssp = ssp + \text{frame size}$)
 - Return by `jr $ra`

Instructions of Procedure Call

- Procedure call: jump and link
 - `jal ProcedureLabel`
 - Address of following instruction put in \$ra
 - Jumps to target address
- Procedure return: jump register
 - `jr $ra`
 - Copies \$ra to program counter
- Can also be used for computed jumps
 - e.g., for case/switch statements

Introduction to Computer Organization _ MIPS Organization

Processor and Memory



Instruction Classes Distribution

Instruction class	MIPS examples	HLL correspondence	Frequency	
			Integer	Ft. pt.
Arithmetic	add, sub, addi	Operations in assignment statement s	16%	48%
Data transfer	lw, sw, lb, lbu, lh, lhu, sb, lui	References to data structures, such as arrays	35%	36%
Logical	and, or, nor, andi, ori, sll, srl	Operations in assignment statement s	12%	4%
Conditional branch	beq, bne, slt, slti, sltiu	If statements and loops	34%	8%
Jump	jr, jal	Procedure calls, returns, and case/switch statements	2%	0%

Addressing

1. Immediate addressing

op rs rt Immediate

2. Register addressing

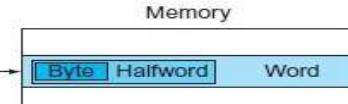
op rs rt rd ... funct

Registers
Register

3. Base addressing

op rs rt Address

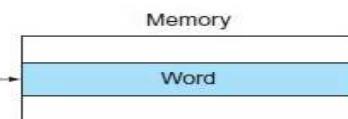
Register



4. PC-relative addressing

op rs rt Address

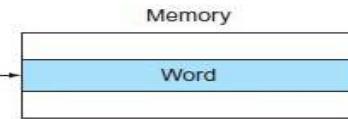
PC



5. Pseudodirect addressing

op Address

PC



MIPS Fields

MIPS fields are given names to make them easier to discuss:

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

Here is the meaning of each name of the fields in MIPS instructions:

- **op:** Basic operation of the instruction, traditionally called the **opcode**.
- **rs:** The first register source operand.
- **rt:** The second register source operand.
- **rd:** The register destination operand. It gets the result of the operation.
- **shamt:** Shift amount. (Section 2.6 explains shift instructions and this term; it will not be used until then, and hence the field contains zero in this section.)
- **funct:** Function. This field, often called the *function code*, selects the specific variant of the operation in the op field.

Introduction to Computer Organization _ MIPS Logic Operations

Logic Operation

Logical operations	C operators	Java operators	MIPS instructions
Shift left	<<	<<	sll
Shift right	>>	>>>	srl
Bit-by-bit AND	&	&	and, andi
Bit-by-bit OR			or, ori
Bit-by-bit NOT	~	~	nor

Basic Operations

- NOT	!A	
- AND	$A \cdot B$	
- OR	$A + B$	
- NOR	$\text{!}(A + B)$	
- NAND	$\text{!}(A \cdot B)$	
- XOR	$A \oplus B$	
- XNOR	$\text{!}(A \oplus B)$	

Basic gates we will use in this class.

AND			OR			XOR		
A	B	Out	A	B	Out	A	B	Out
0	0	0	0	0	0	0	0	0
0	1	0	0	1	1	0	1	1
1	0	0	1	0	1	1	0	1
1	1	1	1	1	1	1	1	0

NAND			NOR			XNOR		
A	B	Out	A	B	Out	A	B	Out
0	0	1	0	0	1	0	0	1
0	1	0	0	1	0	0	1	0
1	0	1	1	0	0	1	0	0
1	1	0	1	1	0	1	1	1

NOT Operation

- Useful to invert bits in a word
 - Change 0 to 1, and 1 to 0
- MIPS has NOR 3-operand instruction
 - $a \text{ NOR } b = \text{NOT} (a \text{ OR } b)$

nor \$t0, \$t1, \$zero

Register \$t0 always treated as zero

\$t1 0000 0000 0000 0000 0011 1100 0000 0000

\$t0 1111 1111 1111 1111 1100 0011 1111 1111

Shift Operation

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- shamt: how many positions to shift
- Shift left logical
 - Shift left and fill with 0 bits
 - sll by i bits multiplies by 2^i
- Shift right logical
 - Shift right and fill with 0 bits
 - srl by i bits divides by 2^i (unsigned only)

AND and OR Operations

- Useful to mask bits in a word
 - Clear some bits to 0, leave others unchanged
 - or \$t0, \$t1, \$t2 # \$t1 is the "AND mask"

\$t2 0000 0000 0000 0000 ... 01 1100 0000

\$t1 0000 0000 0000 0000 ... 00 0000 0000

\$t0 0000 0000 0000 0000 ... 00 0000 0000

- Useful to include bits in a word

$\text{or } \$t0, \$t1, \$t2 \# \$t1 \text{ is the "OR mask"}$

\$t2 0000 0000 0000 0000 ... 01 1100 0000

\$t1 0000 0000 0000 0000 ... 00 1100 0000

\$t0 0000 0000 0000 0000 ... 00 01 1100 0000

Introduction to Computer Organization _ Multiplication Algorithm and Hardware

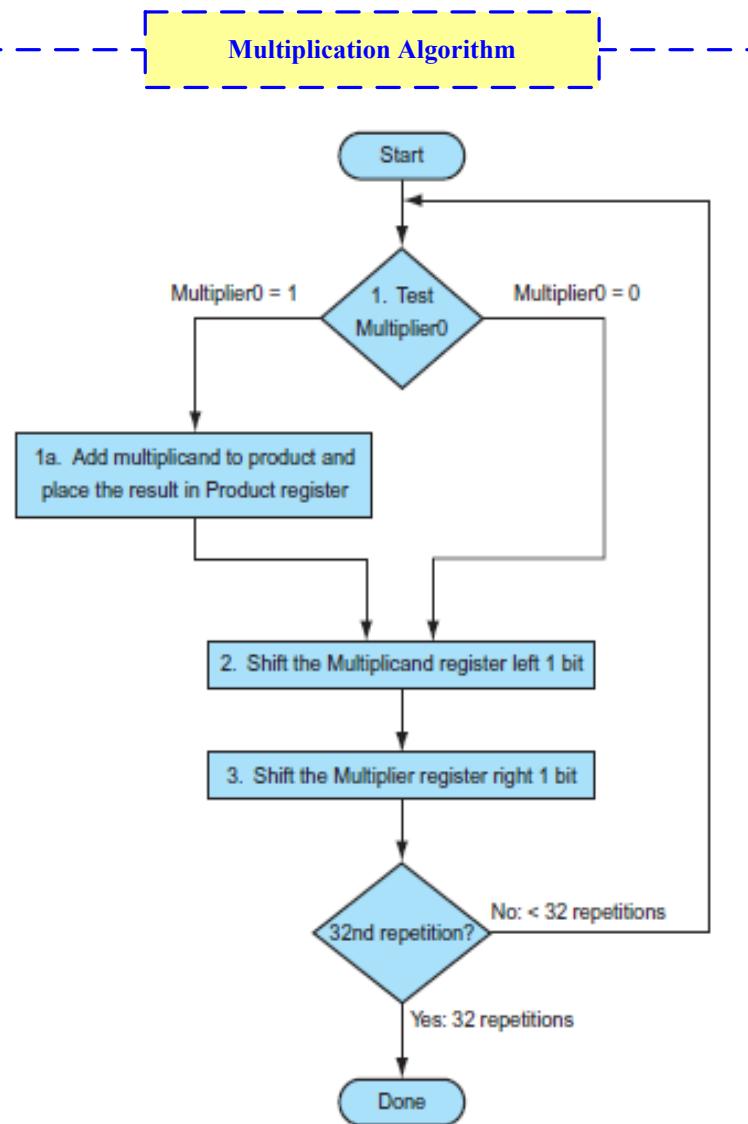


FIGURE 3.4 The first multiplication algorithm, using the hardware shown in Figure 3.3. If the least significant bit of the multiplier is 1, add the multiplicand to the product. If not, go to the next step. Shift the multiplicand left and the multiplier right in the next two steps. These three steps are repeated 32 times.

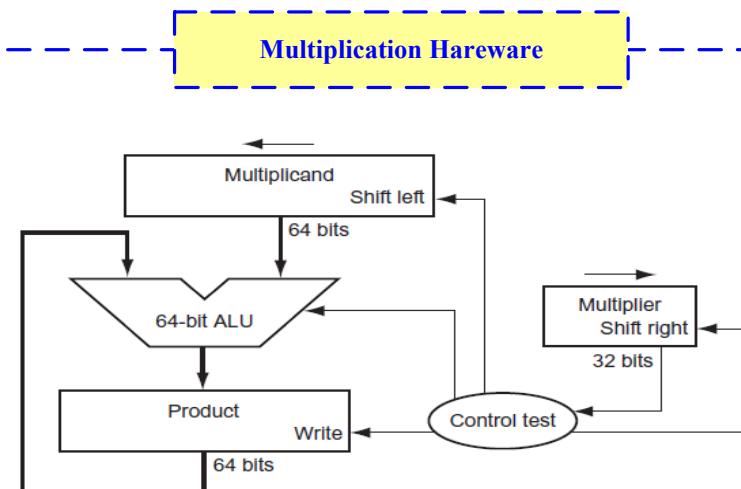


FIGURE 3.3 First version of the multiplication hardware. The Multiplicand register, ALU, and Product register are all 64 bits wide, with only the Multiplier register containing 32 bits. (Appendix B describes ALUs.) The 32-bit multiplicand starts in the right half of the Multiplicand register and is shifted left 1 bit on each step. The multiplier is shifted in the opposite direction at each step. The algorithm starts with the product initialized to 0. Control decides when to shift the Multiplicand and Multiplier registers and when to write new values into the Product register.

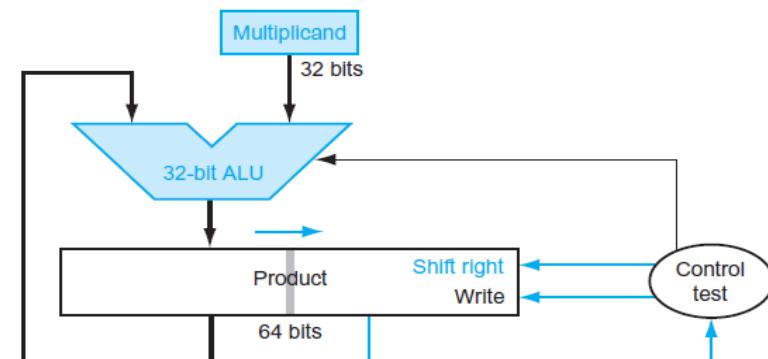


FIGURE 3.5 Refined version of the multiplication hardware. Compare with the first version in Figure 3.3. The Multiplicand register, ALU, and Multiplier register are all 32 bits wide, with only the Product register left at 64 bits. Now the product is shifted right. The separate Multiplier register also disappeared. The multiplier is placed instead in the right half of the Product register. These changes are highlighted in color. (The Product register should really be 65 bits to hold the carry out of the adder, but it's shown here as 64 bits to highlight the evolution from Figure 3.3.)

Introduction to Computer Organization _ Division Algorithm and Hardware

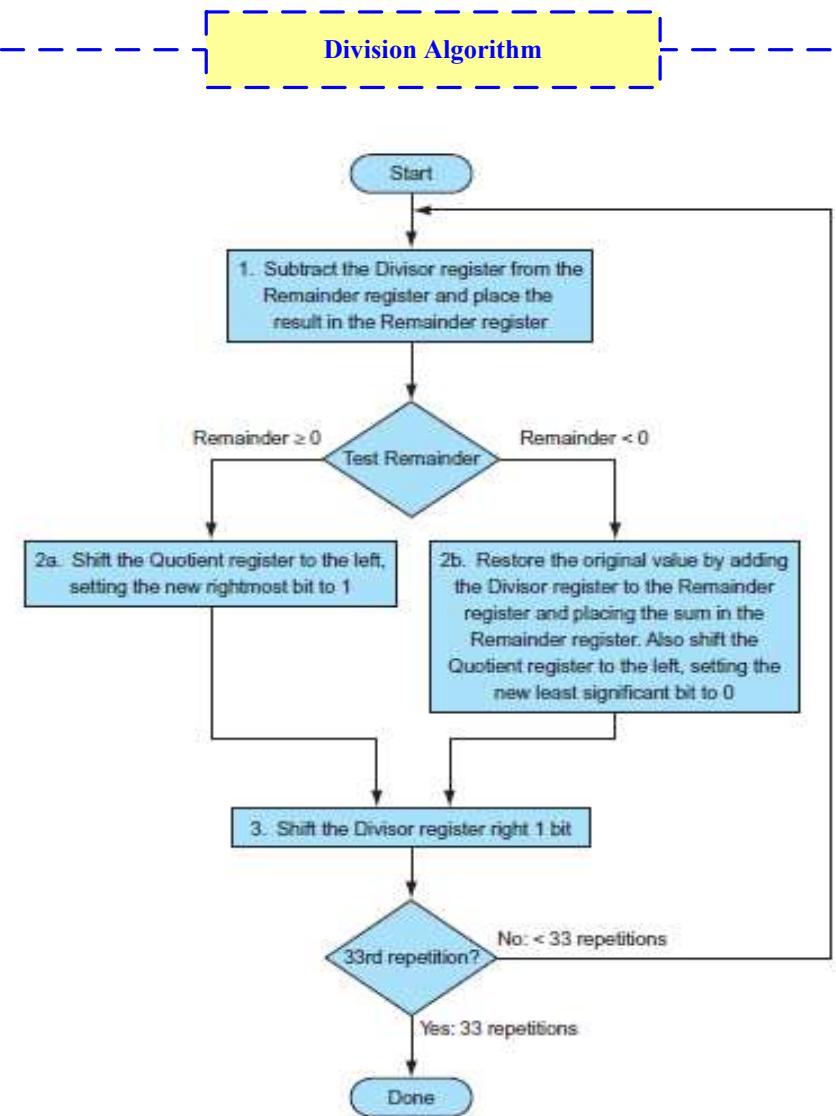


FIGURE 3.9 A division algorithm, using the hardware in Figure 3.8. If the remainder is positive, the divisor did go into the dividend, so step 2a generates a 1 in the quotient. A negative remainder after step 1 means that the divisor did not go into the dividend, so step 2b generates a 0 in the quotient and adds the divisor to the remainder, thereby reversing the subtraction of step 1. The final shift, in step 3, aligns the divisor properly, relative to the dividend for the next iteration. These steps are repeated 33 times.

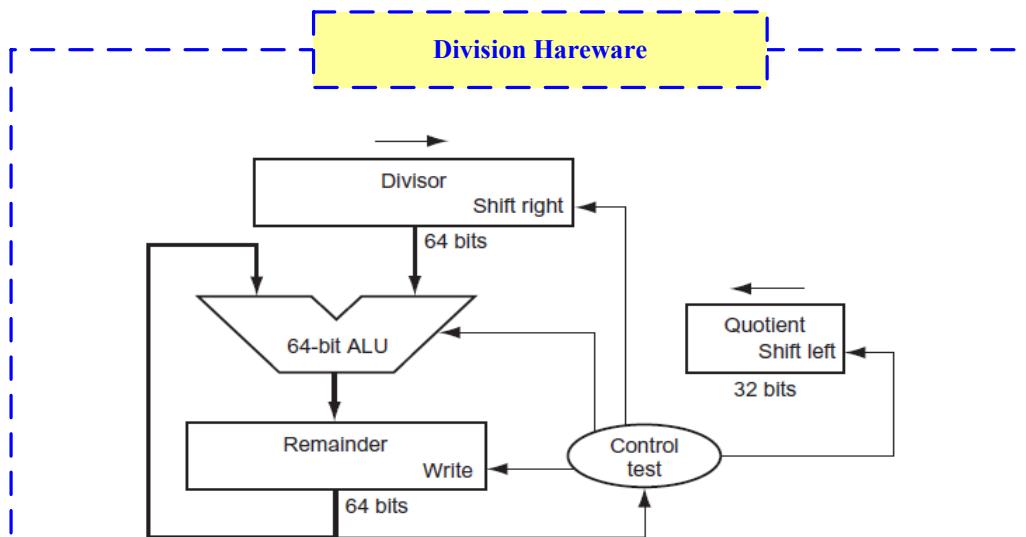


FIGURE 3.8 First version of the division hardware. The Divisor register, ALU, and Remainder register are all 64 bits wide, with only the Quotient register being 32 bits. The 32-bit divisor starts in the left half of the Divisor register and is shifted right 1 bit each iteration. The remainder is initialized with the dividend. Control decides when to shift the Divisor and Quotient registers and when to write the new value into the Remainder register.

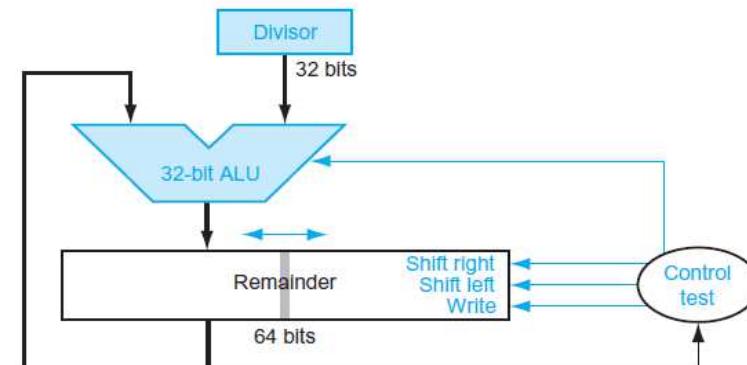


FIGURE 3.11 An improved version of the division hardware. The Divisor register, ALU, and Quotient register are all 32 bits wide, with only the Remainder register left at 64 bits. Compared to Figure 3.8, the ALU and Divisor registers are halved and the remainder is shifted left. This version also combines the Quotient register with the right half of the Remainder register. (As in Figure 3.5, the Remainder register should really be 65 bits to make sure the carry out of the adder is not lost.)

Introduction to Computer Organization _ Floating Point

Floating-Point Representation - single precision

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
s	exponent																														
1 bit	8 bits								23 bits																						

In general, floating-point numbers are of the form

$$(-1)^s \times F \times 2^E$$

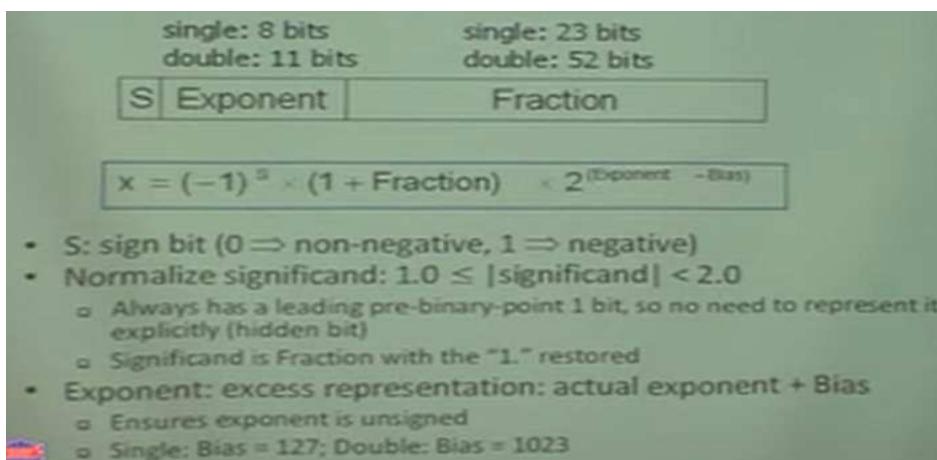
F involves the value in the fraction field and E involves the value in the exponent field; the exact relationship to these fields will be spelled out soon.

Floating-Point Representation - double precision

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
s	exponent																														
1 bit	11 bits											20 bits																			
fraction (continued)																															

MIPS double precision allows numbers almost as small as $2.0_{\text{ten}} \times 10^{-308}$ and almost as large as $2.0_{\text{ten}} \times 10^{308}$

IEEE Floating-Point Format



MIPS Floating-Point

MIPS floating-point operands		
Name	Example	Comments
32 floating-point registers	\$F0, \$F1, \$F2, ..., \$F11	MIPS floating-point registers are used in pairs for double precision numbers.
2 ³² memory words	Memory(\$0), Memory(\$1), ..., Memory(\$4294967292)	Accessed only by data transfer instructions. MIPS uses byte addresses, so sequential word addresses differ by 4. Memory holds data structures, such as arrays, and spilled registers, such as those saved on procedure calls.

MIPS floating-point assembly language

Category	Instruction	Example	Meaning	Comments
Arithmetic	FP add single	add.s \$F2,\$F3,\$F6	\$F2 - FP add + \$F6	FP add (single precision)
	FP subtract single	sub.s \$F2,\$F4,\$F6	\$F2 - FP sub - \$F4 + \$F6	FP sub (single precision)
	FP multiply single	mul.s \$F2,\$F3,\$F6	\$F2 - FP multiply * \$F6	FP multiply (single precision)
	FP divide single	div.s \$F2,\$F3,\$F6	\$F2 - FP divide / \$F6	FP divide (single precision)
	FP add double	add.d \$F2,\$F3,\$F6	\$F2 - FP add + \$F6	FP add (double precision)
	FP subtract double	sub.d \$F2,\$F4,\$F6	\$F2 - FP sub - \$F4 + \$F6	FP sub (double precision)
	FP multiply double	mul.d \$F2,\$F3,\$F6	\$F2 - FP multiply * \$F6	FP multiply (double precision)
	FP divide double	div.d \$F2,\$F3,\$F6	\$F2 - FP divide / \$F6	FP divide (double precision)
Data Transfer	load word coop. t	lwcl \$F1,100(\$s2)	\$F1 - Memory(\$s2 + 100)	32-bit data to FP register
	store word coop. t	swcl \$F1,100(\$s2)	Memory(\$s2 + 100) = \$F1	32-bit data to memory
Conditional Branch	branch on FP true	bclt \$F	If (\$cond == 1) go to PC + 4 + 100	PC-relative branch if FP cond
	branch on FP false	bclf \$F	If (\$cond == 0) go to PC + 4 + 100	PC-relative branch if not cond
	FP compare single	c.lt.s \$F2,\$F4	#(\$F2 < \$F4)	FP compare less than: single precision
	FP compare double	c.lt.d \$F2,\$F4	#(\$F2 < \$F4)	FP compare less than: double precision

MIPS floating-point machine language

Name	Format	Example	Comments	
add.s	R	17	16	6 4 2 0 add.s \$F2,\$F4,\$F6
sub.s	R	17	16	6 4 2 1 sub.s \$F2,\$F4,\$F6
mul.s	R	17	16	6 4 2 3 mul.s \$F2,\$F4,\$F6
div.s	R	17	16	6 4 2 3 div.s \$F2,\$F4,\$F6
add.d	R	17	17	6 4 2 0 add.d \$F2,\$F4,\$F6
sub.d	R	17	17	6 4 2 1 sub.d \$F2,\$F4,\$F6
mul.d	R	17	17	6 4 2 2 mul.d \$F2,\$F4,\$F6
div.d	R	17	17	6 4 2 3 div.d \$F2,\$F4,\$F6
lwcl	I	49	20	2 100 lwcl \$F2,100(\$s4)
swcl	I	57	20	2 100 swcl \$F2,100(\$s4)
bclt	I	17	B	1 25 bclt 25
bclf	I	17	B	0 25 bclf 25
c.lt.s	R	17	16	4 2 0 60 c.lt.s \$F2,\$F4
c.lt.d	R	17	17	4 2 0 60 c.lt.d \$F2,\$F4
Field size		6 bits	5 bits	5 bits 5 bits 6 bits All MIPS instructions 32 bits

FIGURE 3.17 MIPS floating-point architecture revealed thus far. See Appendix A, Section A.10, for more detail. This information is also found in column 2 of the MIPS Reference Data Card at the front of this book.

Introduction to Computer Organization _ Floating-Point Addition

Floating-Point - Multiplication Algorithm

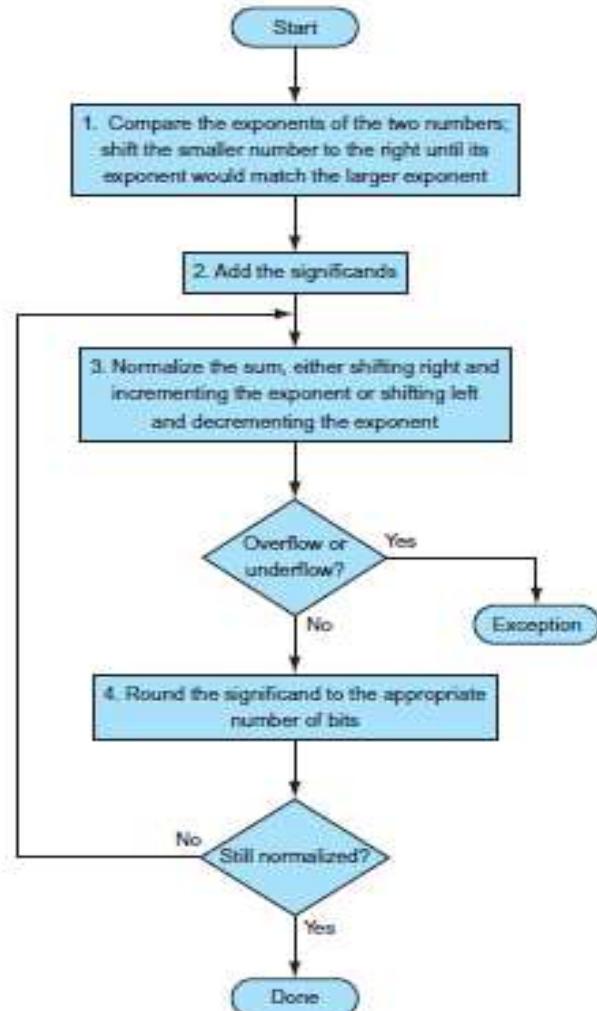


FIGURE 3.14 Floating-point addition. The normal path is to execute steps 3 and 4 once, but if rounding causes the sum to be unnormalized, we must repeat step 3.

Floating-Point - Addition Hardware

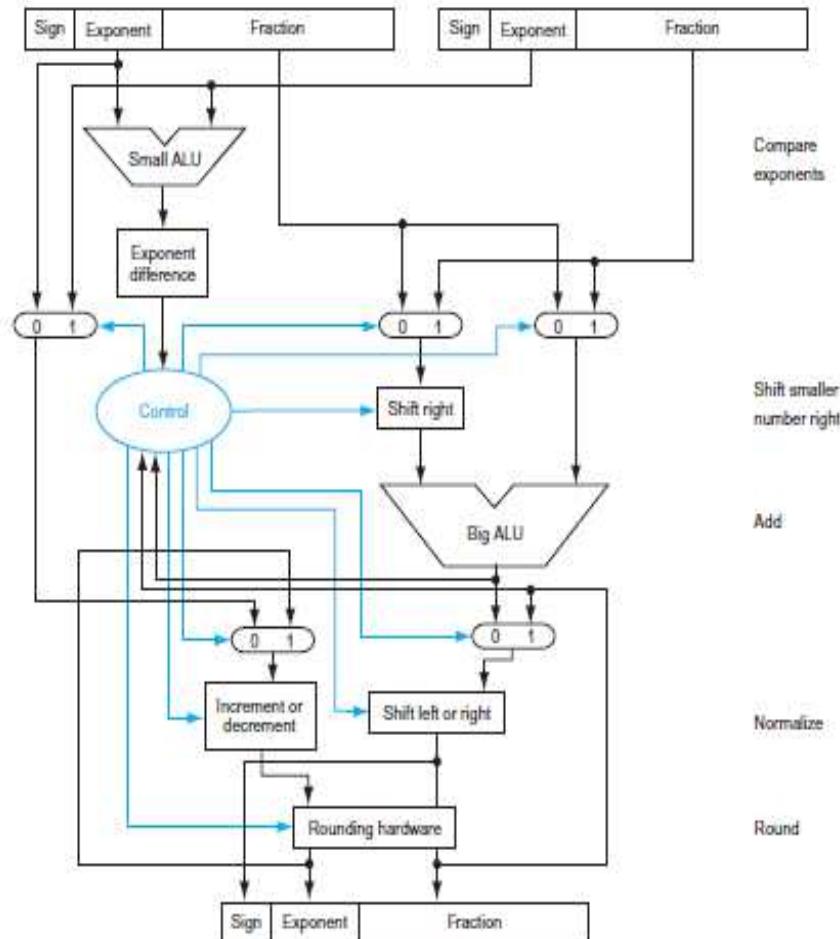
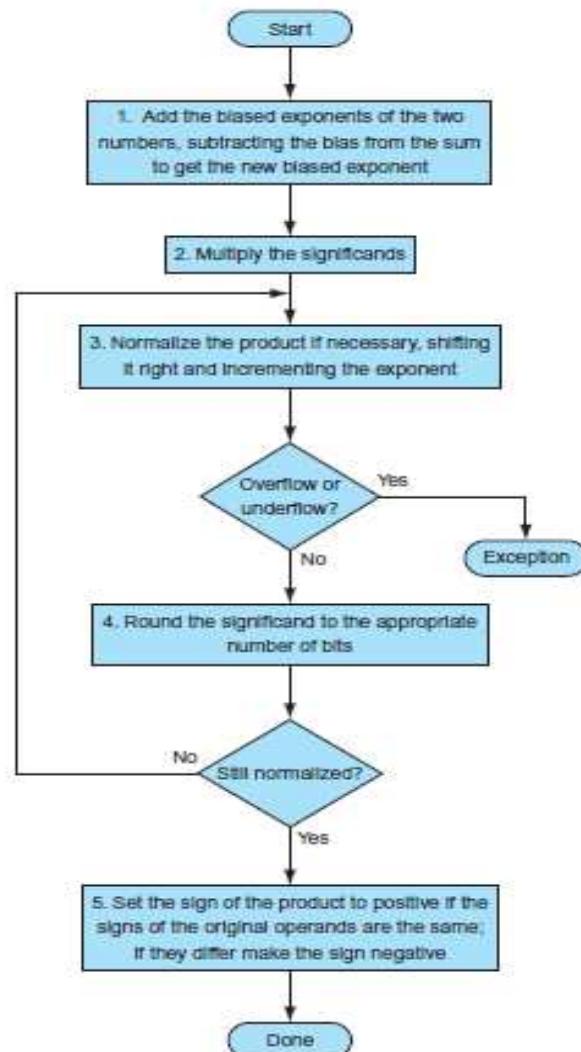


FIGURE 3.15 Block diagram of an arithmetic unit dedicated to floating-point addition. The steps of Figure 3.14 correspond to each block, from top to bottom. First, the exponent of one operand is subtracted from the other using the small ALU to determine which is larger and by how much. This difference controls the three multiplexors from left to right, they select the larger exponent, the significand of the smaller number, and the significand of the larger number. The smaller significand is shifted right, and then the significands are added together using the big ALU. The normalization step then shifts the sum left or right and increments or decrements the exponent. Rounding then creates the final result, which may require normalizing again to produce the actual final result.

Introduction to Computer Organization _ Floating-Point Multiplication

Floating-Point - Multiplication Algorithm



Floating-Point Summary

The major difference between computer numbers and numbers in the real world is that computer numbers have limited size and hence limited precision; it's possible to calculate a number too big or too small to be represented in a word. Programmers must remember these limits and write programs accordingly.

C type	Java type	Data transfers	Operations
int	int	lw, sw, lui	addu, addiu, subu, mult, div, AND, ANDi, OR, ORi, NOR, slt, slti
unsigned int	-	lw, sw, lui	addu, addiu, subu, multu, divu, AND, ANDi, OR, ORi, NOR, sltu, sltiu
char	-	lb, sb, lui	add, addi, sub, mult, div AND, ANDi, OR, ORi, NOR, slt, slti
-	char	lh, sh, lui	addu, addiu, subu, multu, divu, AND, ANDi, OR, ORi, NOR, sltu, sltiu
float	float	lwcl, swcl	add.s, sub.s, mult.s, div.s, c.eq.s, c.lt.s, c.le.s
double	double	ld, sd	add.d, sub.d, mult.d, div.d, c.eq.d, c.lt.d, c.le.d

FIGURE 3.16 Floating-point multiplication. The normal path is to execute steps 3 and 4 once, but if rounding causes the sum to be unnormalized, we must repeat step 3.

Introduction to Computer Organization _ Performance Thinking

Performance Big View

- ❑ Measure, Report, and Summarize
- ❑ Make intelligent choices
- ❑ See through the marketing hype
- ❑ Key to understanding underlying organizational motivation

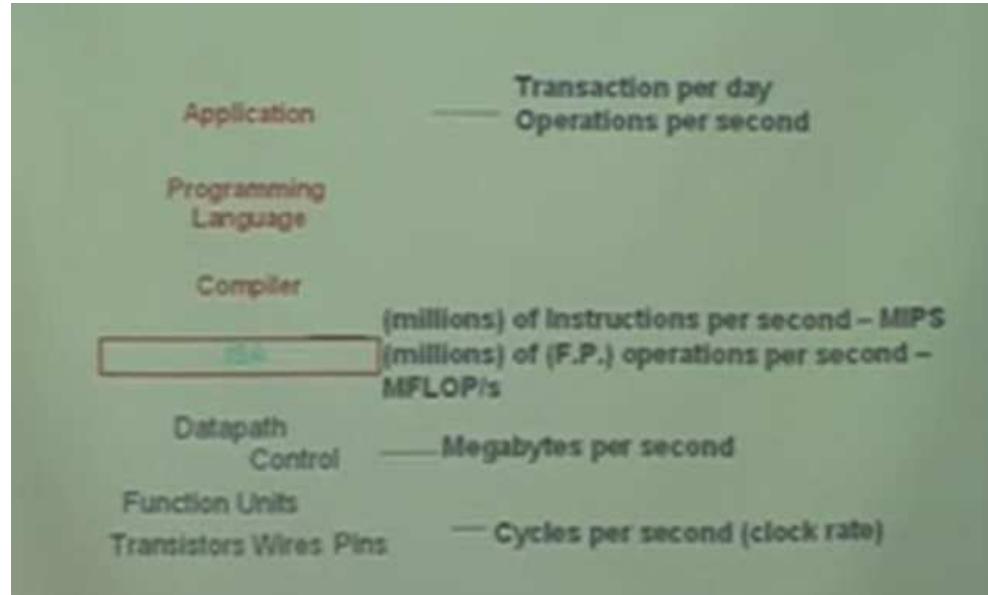
Why is some hardware better than others for different programs?

*What factors of system performance are hardware related?
(e.g., Do we need a new machine, or a new operating system?)*

- ❑ Response Time (latency)
 - How long does it take for my job to run?
 - How long does it take to execute a job?
 - How long must I wait for the database query?
- ❑ Throughput
 - How many jobs can the machine run at once?
 - What is the average execution rate?
 - How much work is getting done?
- ❑ If we upgrade a machine with a new processor what do we increase?
- ❑ If we add a new machine to the lab what do we increase?

Performance Metrics

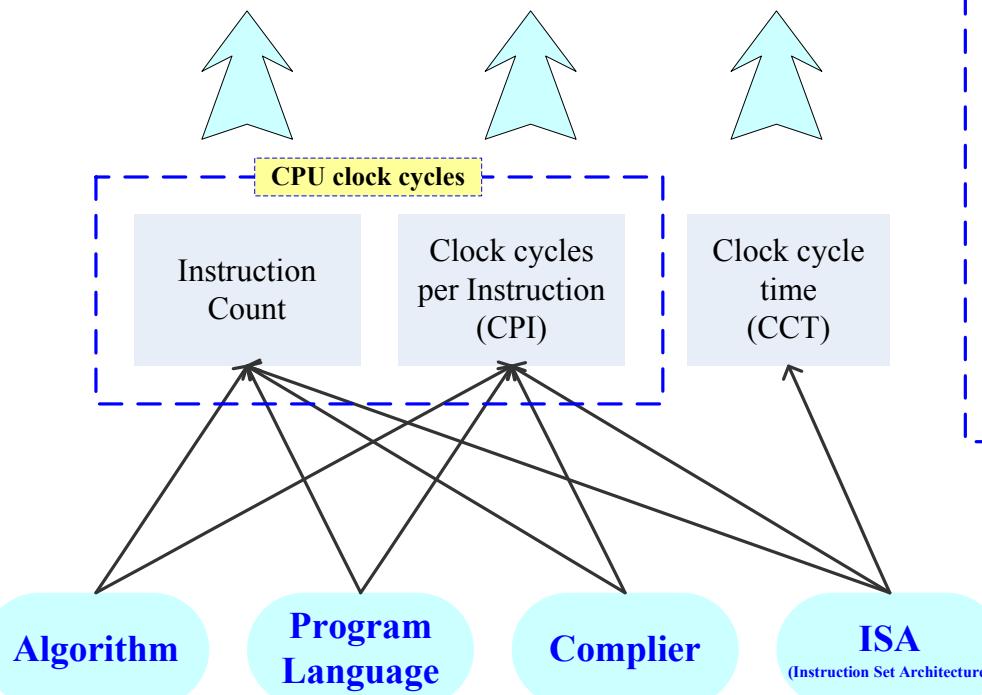
- ❑ Purchasing perspective
 - given a collection of machines, which has the
 - best performance ?
 - least cost ?
 - best cost/performance?
- ❑ Design perspective
 - faced with design options, which has the
 - best performance improvement ?
 - least cost ?
 - best cost/performance?
- ❑ Both require
 - basis for comparison
 - metric for evaluation
- ❑ Our goal is to understand what factors in the architecture contribute to overall system performance and the relative importance (and cost) of these factors



Introduction to Computer Organization _ Performance

Performance

$$\text{Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$



MIPS

(million instructions per second)

$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}$$

$$\text{MIPS} = \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}} \times 10^6} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

Introduction to Computer Organization _ Processor Design _ Logic Unit

Logic Operation

Basic Operations

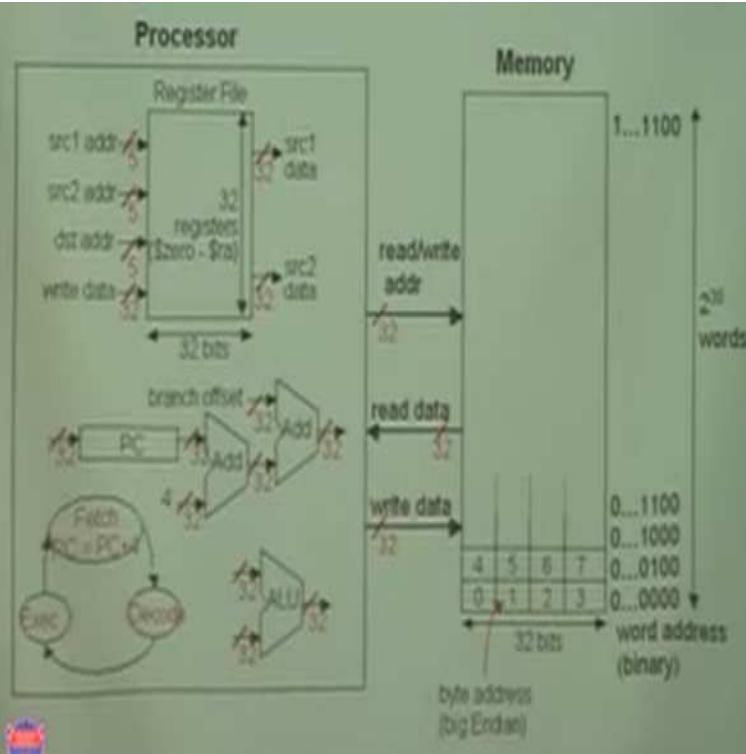
- NOT	$\neg A$
- AND	$A \cdot B$
- OR	$A + B$
- NOR	$\neg(A + B) = \neg(A \text{ OR } B)$
- NAND	$\neg(A \cdot B) = \neg(A \text{ AND } B)$
- XOR	$A \oplus B = (A \cdot \neg B) + (\neg A \cdot B)$
- XNOR	$\neg(A \oplus B) = \neg(A \text{ XOR } B)$

Basic gates we will use in this class

AND			OR			XOR		
A	B	Out	A	B	Out	A	B	Out
0	0	0	0	0	0	0	0	0
0	1	0	0	1	1	0	1	1
1	0	0	1	0	1	1	0	1
1	1	1	1	1	1	1	1	0

NAND			NOR			XNOR		
A	B	Out	A	B	Out	A	B	Out
0	0	1	0	0	1	0	0	1
0	1	1	0	1	0	0	1	0
1	0	1	1	0	0	1	0	0
1	1	0	1	1	0	1	1	1

Logic Blocks

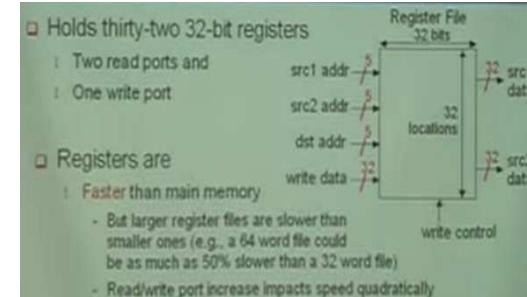


Processor

Memory

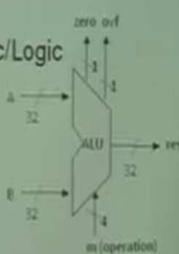
Logic Blocks

- MUXes** select an input from many
 - The input may be a bus (multiple bits)
- DEMUXes** do the opposite: output to one output
- DECODERS** take a binary valued input and activate one output (1-hot)
 - E.g., binary in 010 will activate 1-hot output #2
- ENCODERS** take a 1-hot input and output its binary value
 - E.g., 1-hot input #3 in will output binary 011
- ADDERS** take in A and B and output the sum and the carry
 - Sum is the same number of bits as A or B
 - Carry is one additional bit in case you have an overflow



ALU

- Must support the Arithmetic/Logic operations of the ISA
- add, addi, addiu, addu
- sub, subu
- mult, multu, div, divu
- sqrt
- and, andi, nor, or, ori, xor, xorl
- beq, bne,slt, slti, sltiu, sltiu

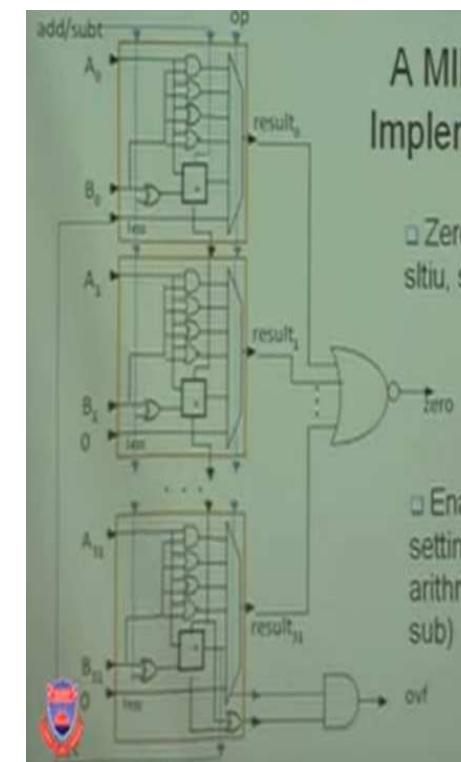


Operation	Operand A	Operand B	Result indicating overflow
$A + B$	≥ 0	≥ 0	< 0
$A + B$	< 0	< 0	≥ 0
$A - B$	≥ 0	≥ 0	< 0
$A - B$	< 0	≥ 0	≥ 0

FIGURE 3.2 Overflow conditions for addition and subtraction.

A MIPS ALU Implementation

- Zero detect (slt, slti, sltiu, sltu, beq, bne)



- Enable overflow bit setting for signed arithmetic (add, addi, sub)

Introduction to Computer Organization _ Processor Design _ Datapath

Elements of Datapath

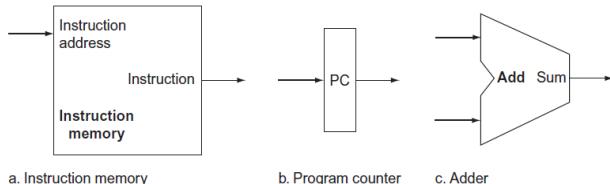


FIGURE 4.5 Two state elements are needed to store and access instructions, and an adder is needed to compute the next instruction address.

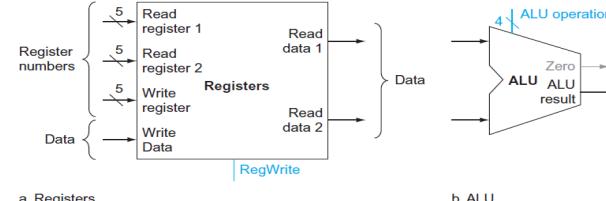


FIGURE 4.7 The two elements needed to implement R-format ALU operations are the register file and the ALU.

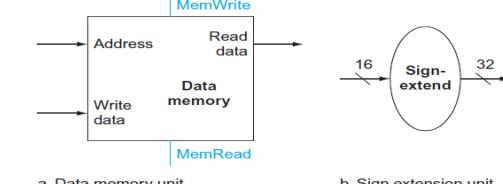


FIGURE 4.8 The two units needed to implement loads and stores, in addition to the register file and ALU of Figure 4.7, are the data memory unit and the sign extension unit.

Design of Datapath

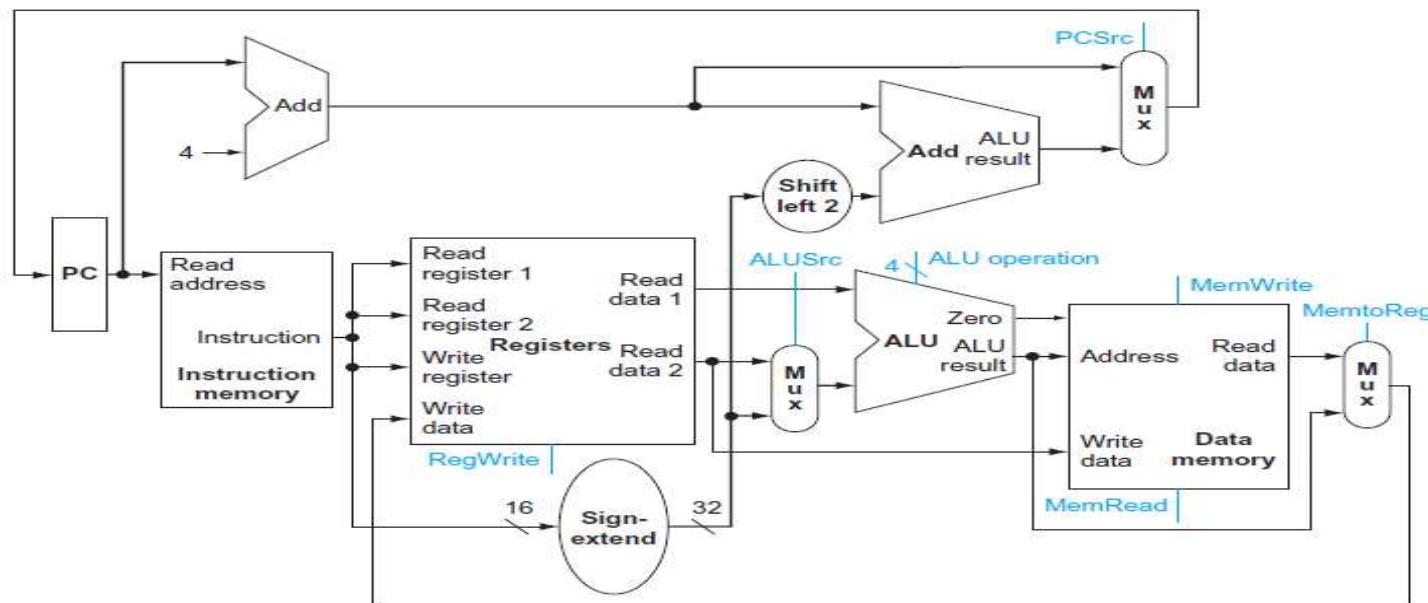


FIGURE 4.11 The simple datapath for the core MIPS architecture combines the elements required by different instruction classes. The components come from Figures 4.6, 4.9, and 4.10. This datapath can execute the basic instructions (load-store word, ALU operations, and branches) in a single clock cycle. Just one additional multiplexor is needed to integrate branches. The support for jumps will be added later.

Introduction to Computer Organization _ Processor Design _ Control

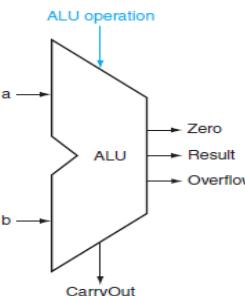


FIGURE B.5.14 The symbol commonly used to represent an ALU, as shown in Figure B.5.12. This symbol is also used to represent an adder, so it is normally labeled either with ALU or Adder.

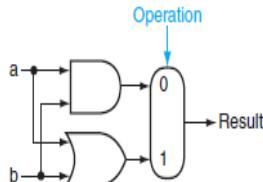


FIGURE B.5.1 The 1-bit logical unit for AND and OR.

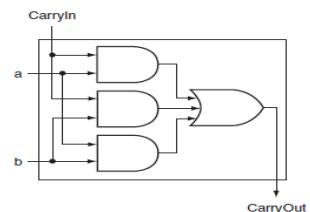


FIGURE B.5.5 Adder hardware for the CarryOut signal

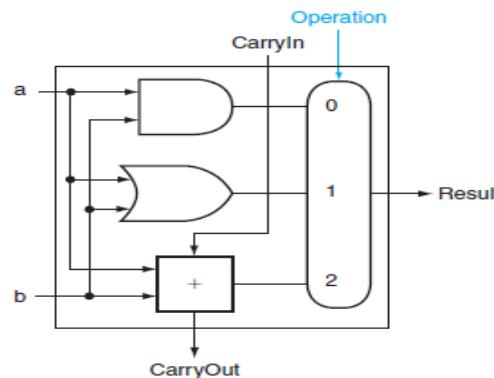


FIGURE B.5.6 A 1-bit ALU that performs AND, OR, and addition (see Figure B.5.5).

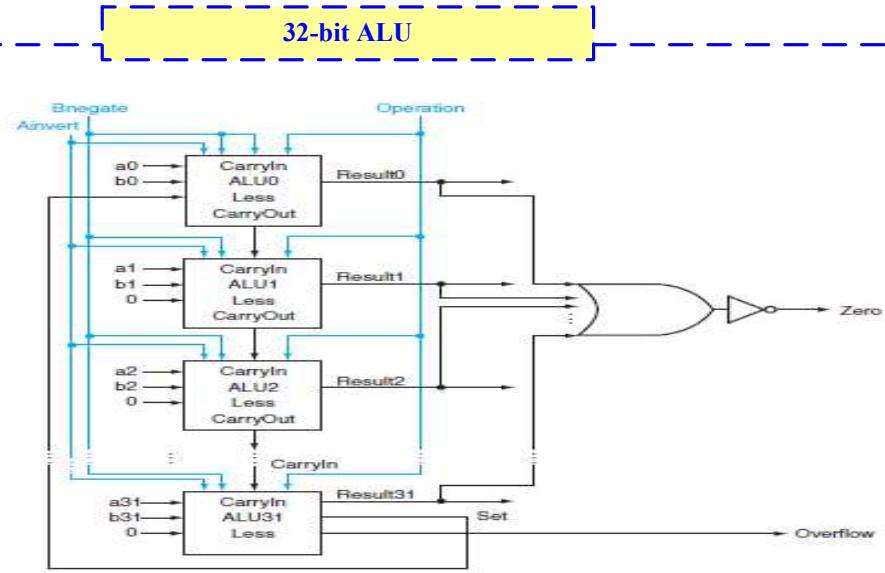


FIGURE B.5.12 The final 32-bit ALU. This adds a Zero detector to Figure B.5.11.

Instruction opcode	ALUOp	Instruction operation	Funct field	Desired ALU action	ALU control input
LW	00	load word	XXXXXX	add	0010
SW	00	store word	XXXXXX	add	0010
Branch equal	01	branch equal	XXXXXX	subtract	0110
R-type	10	add	100000	add	0010
R-type	10	subtract	100010	subtract	0110
R-type	10	AND	100100	AND	0000
R-type	10	OR	100101	OR	0001
R-type	10	set on less than	101010	set on less than	0111

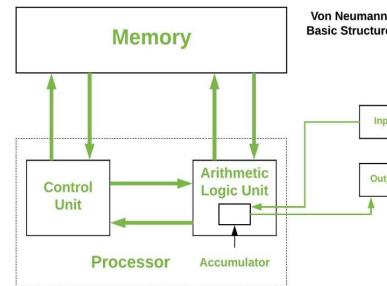
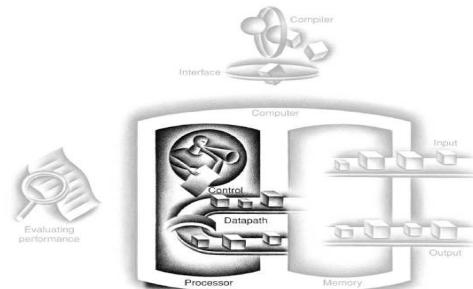
FIGURE 4.12 How the ALU control bits are set depends on the ALUOp control bits and the different function codes for the R-type instruction. See Appendix B.

ALUOp		Funct field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
0	0	X	X	X	X	X	X	0010
X	1	X	X	X	X	X	X	0110
1	X	X	X	0	0	0	0	0010
1	X	X	X	0	0	1	0	0110
1	X	X	X	0	1	0	0	0000
1	X	X	X	0	1	0	1	0001
1	X	X	X	1	0	1	0	0111

FIGURE 4.13 The truth table for the 4 ALU control bits (called Operation). The inputs are the ALUOp and function code field. Only the entries for which the ALU control is asserted are shown. Some don't-care entries have been added. For example, the ALUOp does not use the encoding 11, so the truth table can contain entries 1X and X1, rather than 10 and 01. Note that when the function field is used, the first 2 bits (F5 and F4) of these instructions are always 10, so they are don't-care terms and are replaced with XX in the truth table.

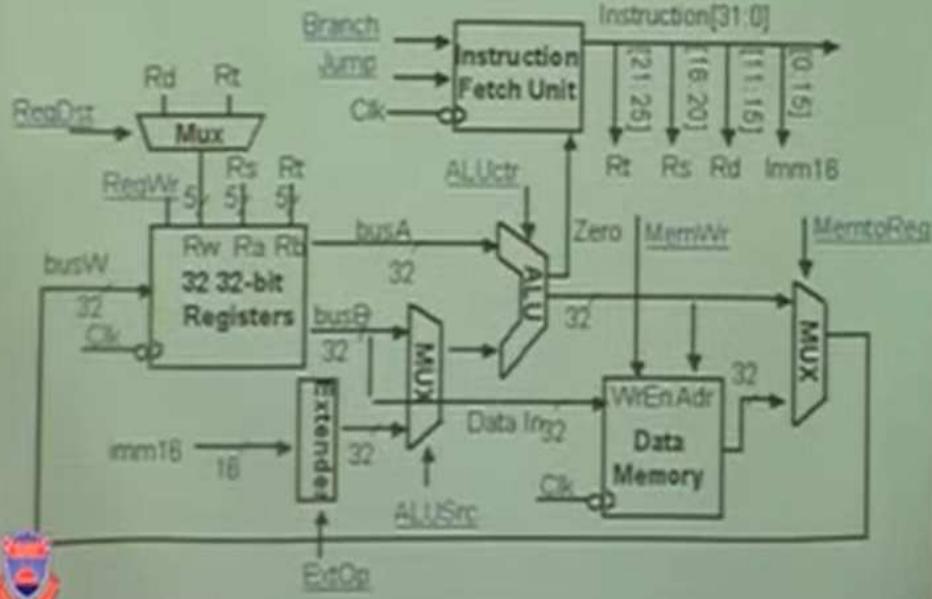
Introduction to Computer Organization _ Processor Design _ ALL

Five Classic Components of a Computer



Put it together

Putting it All Together

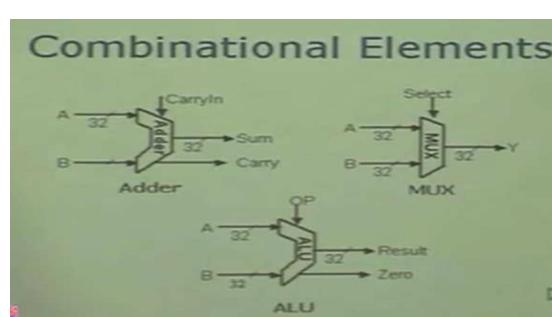


Processor Design Steps

- Analyze instruction set => datapath requirements**
 - the meaning of each instruction is given by the register transfers (ISA model \rightarrow RTL model)
 - datapath must include storage element for ISA registers
 - possibly more
 - datapath must support each register transfer
- Select set of datapath components and establish clocking methodology**
- Assemble datapath meeting the RTL requirements**
- Analyze implementation of each instruction to determine setting of control points that effect the assembly**
- Assemble the control logic**
- RTL datapath and control design are refined to track physical design and functional validation**
 - Changes made for timing and errata (a.k.a. "bug") fixes
 - Amount of work varies with capabilities of CAD tools and degree of optimization for cost/performance

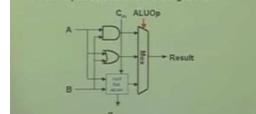
Logic Design

Combinational Elements

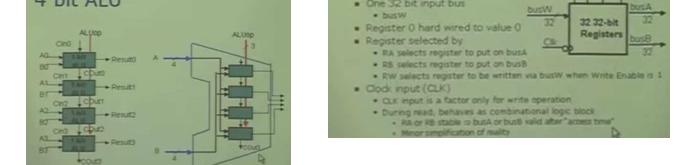


1 bit ALU

- Using a MUX we can add the AND, OR, and adder operations into a single ALU



4 bit ALU



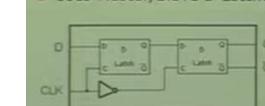
D Latches

- Modified SR Latch
- Latches value when C is asserted

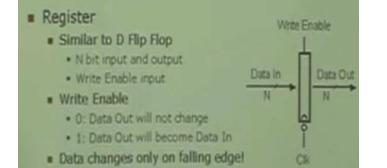


D Flip Flops

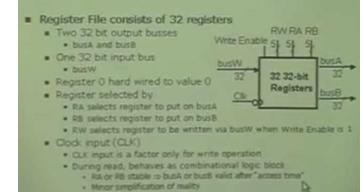
- Uses Master/Slave D Latches



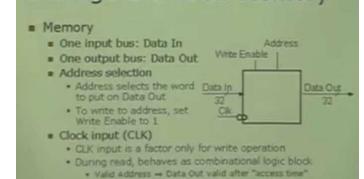
Storage Element: Register



Storage Element: Reg File

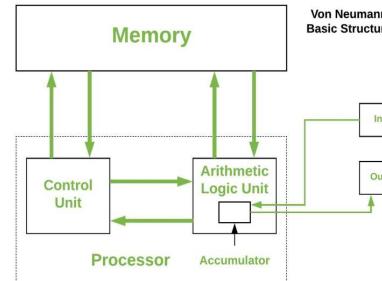
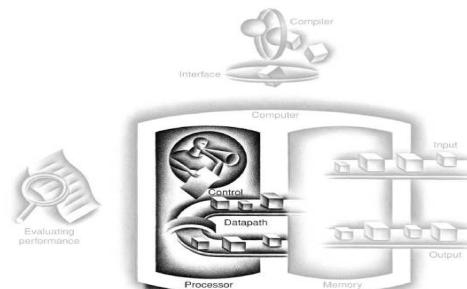


Storage Element: Memory



Introduction to Computer Organization _ Processor Design _ ALL

Five Classic Components of a Computer



Operand of the Instruction

Instruction	RegDst	ALUSrc	MemtoReg	RegWrite	MemRead	MemWrite	Branch	ALUOp1	ALUOp0
R-format	1	0	0	1	0	0	0	1	0
lw	0	1	1	1	1	0	0	0	0
sw	X	1	X	0	0	1	0	0	0
beq	X	0	X	0	0	0	1	0	1

FIGURE 4.18 The setting of the control lines is completely determined by the opcode fields of the instruction. The first row of the table corresponds to the R-format instructions (add, sub, AND, OR, and slt). For all these instructions, the source register fields are rs and rt, and the destination register field is rd; this defines how the signals ALUSrc and RegDst are set. Furthermore, an R-type instruction writes a register (RegWrite = 1), but neither reads nor writes data memory. When the Branch control signal is 0, the PC is unconditionally replaced with PC + 4; otherwise, the PC is replaced by the branch target if the Zero output of the ALU is also high. The ALUOp field for R-type instructions is set to 10 to indicate that the ALU control should be generated from the funct field. The second and third rows of this table give the control signal settings for lw and sw. These ALUSrc and ALUOp fields are set to perform the address calculation. The MemRead and MemWrite are set to perform the memory access. Finally, RegDst and RegWrite are set for a load to cause the result to be stored into the rt register. The branch instruction is similar to an R-format operation, since it sends the rs and rt registers to the ALU. The ALUOp field for branch is set for a subtract (ALU control = 01), which is used to test for equality. Notice that the MemtoReg field is irrelevant when the RegWrite signal is 0; since the register is not being written, the value of the data on the register data write port is not used. Thus, the entry MemtoReg in the last two rows of the table is replaced with X for don't care. Don't cares can also be added to RegDst when RegWrite is 0. This type of don't care must be added by the designer, since it depends on knowledge of how the datapath works.

The Implementation of the MIPS subset

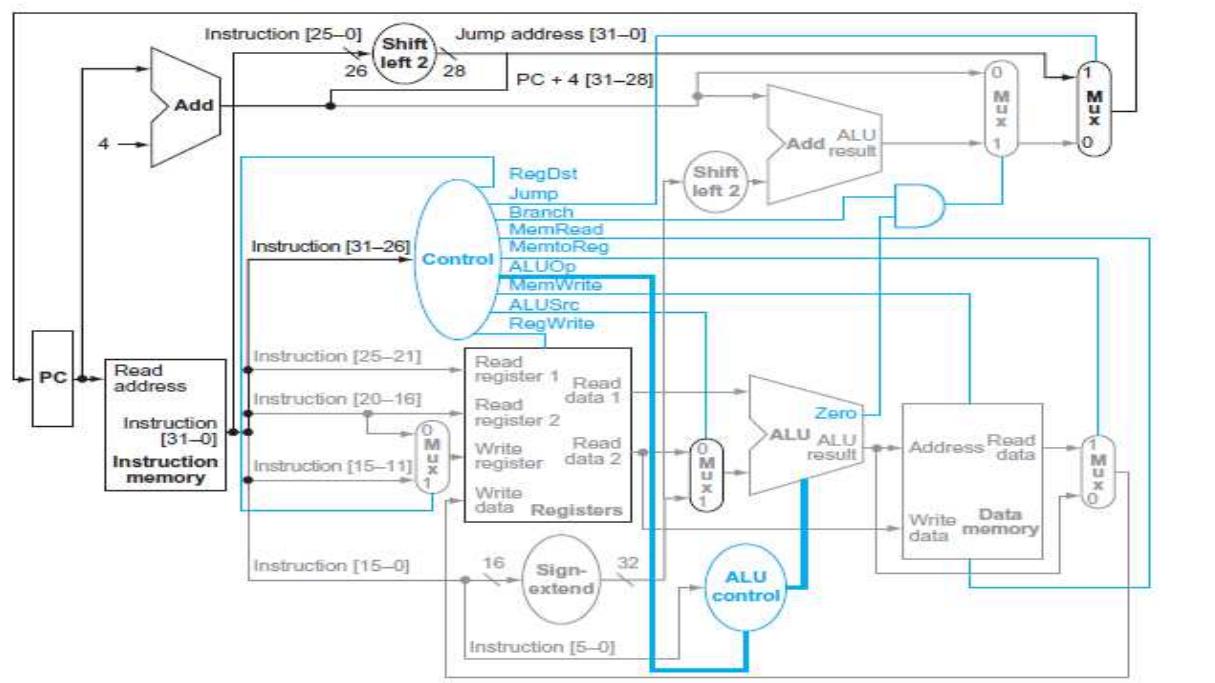


FIGURE 4.24 The simple control and datapath are extended to handle the jump instruction.

Signal name	Effect when deasserted	Effect when asserted
RegDst	The register destination number for the Write register comes from the rt field (bits 20:16).	The register destination number for the Write register comes from the rd field (bits 15:11).
RegWrite	None.	The register on the Write register input is written with the value on the Write data input.
ALUSrc	The second ALU operand comes from the second register file output (Read data 2).	The second ALU operand is the sign-extended, lower 16 bits of the instruction.
PCSrc	The PC is replaced by the output of the adder that computes the value of PC + 4.	The PC is replaced by the output of the adder that computes the branch target.
MemRead	None.	Data memory contents designated by the address input are put on the Read data output.
MemWrite	None.	Data memory contents designated by the address input are replaced by the value on the Write data input.
MemtoReg	The value fed to the register Write data input comes from the data memory.	The value fed to the register Write data input comes from the ALU.

FIGURE 4.16 The effect of each of the seven control signals.

Input or output	Signal name	R-format	lw	sw	beq
Inputs	Op5	0	1	1	0
	Op4	0	0	0	0
	Op3	0	0	1	0
	Op2	0	0	0	1
	Op1	0	1	1	0
	Op0	0	1	1	0
Outputs	RegDst	1	0	X	X
	ALUSrc	0	1	1	0
	MemtoReg	0	1	X	X
	RegWrite	1	1	0	0
	MemRead	0	1	0	0
	MemWrite	0	0	1	0
	Branch	0	0	0	1
	ALUOp1	1	0	0	0
	ALUOp0	0	0	0	1

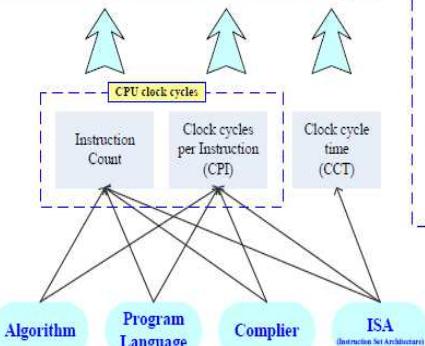
FIGURE 4.22 The control function for the simple single-cycle implementation is completely specified by this truth table.

Introduction to Computer Organization _ Processor Design _ Pipeline _ Basic

Why?

Performance

$$\text{Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$



MIPS (million instructions per second)

$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time}} \times 10^6$$

$$\text{MIPS} = \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}} \times 10^6} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

MIPS Pipeline 5 Stages

MIPS instructions classically take five steps:

1. Fetch instruction from memory.
2. Read registers while decoding the instruction. The regular format of MIPS instructions allows reading and decoding to occur simultaneously.
3. Execute the operation or calculate an address.
4. Access an operand in data memory.
5. Write the result into a register.

Instruction class	Instruction fetch	Register read	ALU operation	Data access	Register write	Total time
Load word (lw)	200 ps	100 ps	200 ps	200 ps	100 ps	800 ps
Store word (sw)	200 ps	100 ps	200 ps	200 ps		700 ps
R-format (add, sub, AND, OR, slt)	200 ps	100 ps	200 ps		100 ps	600 ps
Branch (beq)	200 ps	100 ps	200 ps			500 ps

FIGURE 4.26 Total time for each instruction calculated from the time for each component.
This calculation assumes that the multiplexors, control unit, PC accesses, and sign extension unit have no delay.

Pipeline Hazards

structural hazard

When a planned instruction cannot execute in the proper clock cycle because the hardware does not support the combination of instructions that are set to execute.

data hazard

When a planned instruction cannot execute in the proper clock cycle because data that is needed to execute the instruction is not yet available.

control hazard

When the proper instruction cannot execute in the proper pipeline clock cycle because the instruction that was fetched is not the one that is needed; that is, the flow of instruction addresses is not what the pipeline expected.

Introduction to Computer Organization _ Processor Design _ Pipeline _ Control

Pipeline Design

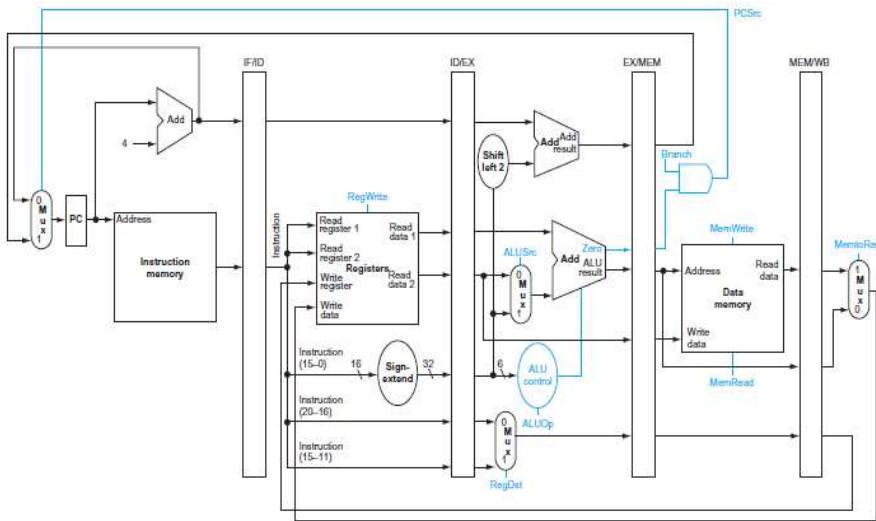


FIGURE 4.46 The pipelined datapath of Figure 4.41 with the control signals identified. This datapath borrows the control logic for PC source, register destination number, and ALU control from Section 4.4. Note that we now need the 6-bit funct field (function code) of the instruction in the EX stage as input to ALU control, so these bits must also be included in the ID/EX pipeline register. Recall that these 6 bits are also the 6 least significant bits of the immediate field in the instruction, so the ID/EX pipeline register can supply them from the immediate field since sign extension leaves these bits unchanged.

32-bit ALU

Instruction opcode	ALUOp	Instruction operation	Function code	Desired ALU action	ALU control input
LW	00	load word	XXXXXX	add	0010
SW	00	store word	XXXXXX	add	0010
Branch equal	01	branch equal	XXXXXX	subtract	0110
R-type	10	add	100000	add	0010
R-type	10	subtract	100010	subtract	0110
R-type	10	AND	100100	AND	0000
R-type	10	OR	100101	OR	0001
R-type	10	set on less than	101010	set on less than	0111

FIGURE 4.47 A copy of Figure 4.12. This figure shows how the ALU control bits are set depending on the ALUOp control bits and the different function codes for the R-type instruction.

Signal name	Effect when deasserted (0)	Effect when asserted (1)
RegDst	The register destination number for the Write register comes from the rt field (bits 20:16).	The register destination number for the Write register comes from the rd field (bits 15:11).
RegWrite	None.	The register on the Write register input is written with the value on the Write data input.
ALUSrc	The second ALU operand comes from the second register file output (Read data 2).	The second ALU operand is the sign-extended, lower 16 bits of the instruction.
PCSrc	The PC is replaced by the output of the adder that computes the branch target.	The PC is replaced by the output of the adder that computes the branch target.
MemRead	None.	Data memory contents designated by the address input are put on the Read data output.
MemWrite	None.	Data memory contents designated by the address input are replaced by the value on the Write data input.
MemtoReg	The value fed to the register Write data input	The value fed to the register Write data input comes from the data memory.

FIGURE 4.48 A copy of Figure 4.16. The function of each of seven control signals is defined. The ALU control lines (ALUOp) are defined in the second column of Figure 4.47. When a 1-bit control to a 2-way multiplexer is asserted, the multiplexer selects the input corresponding to 1. Otherwise, if the control is deasserted, the multiplexer selects the 0 input. Note that PCSrc is controlled by an AND gate in Figure 4.46. If the Branch signal and the ALU Zero signal are both set, then PCSrc is 1; otherwise, it is 0. Control sets the Branch signal only during a beq instruction; otherwise, PCSrc is set to 0.

final three stages

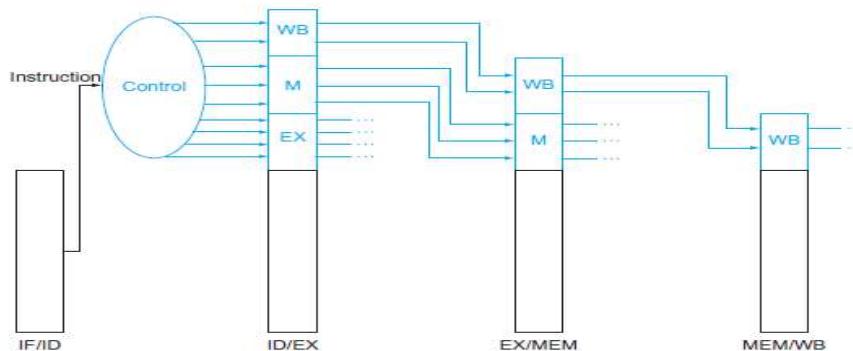


FIGURE 4.50 The control lines for the final three stages. Note that four of the nine control lines are used in the EX phase, with the remaining five control lines passed on to the EX/MEM pipeline register extended to hold the control lines; three are used during the MEM stage, and the last two are passed to MEM/WB for use in the WB stage.

Execution/address calculation stage control lines

Memory access stage control lines

Write-back stage control lines

Instruction	Execution/address calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	RegDst	ALUOp1	ALUOp0	ALUSrc	Branch	Mem-Read	Mem-Write	Reg-Write	Memto-Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

FIGURE 4.49 The values of the control lines are the same as in Figure 4.18, but they have been shuffled into three groups corresponding to the last three pipeline stages.

Introduction to Computer Organization _ Processor Design _ Pipeline _ Datapath

Pipeline Design

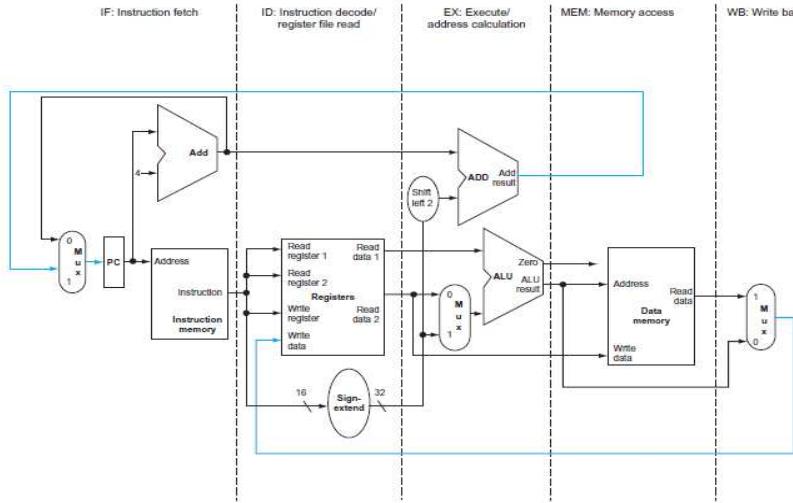


FIGURE 4.33 The single-cycle datapath from Section 4.4 (similar to Figure 4.17). Each step of the instruction can be mapped onto the datapath from left to right. The only exceptions are the update of the PC and the write-back step, shown in color, which sends either the ALU result or the data from memory to the left to be written into the register file. (Normally we use color lines for control, but these are data lines.)

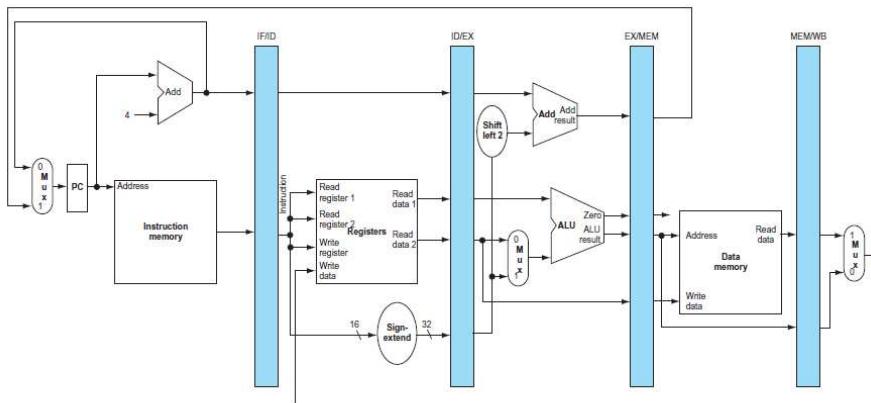


FIGURE 4.35 The pipelined version of the datapath in Figure 4.33. The pipeline registers, in color, separate each pipeline stage. They are labeled by the stages that they separate; for example, the first is labeled IF/ID because it separates the instruction fetch and instruction decode stages. The registers must be wide enough to store all the data corresponding to the lines that go through them. For example, the IF/ID register must be 64 bits wide, because it must hold both the 32-bit instruction fetched from memory and the incremented 32-bit PC address. We will expand these registers over the course of this chapter, but for now the other three pipeline registers contain 128, 97, and 64 bits, respectively.

Pipeline Datapath

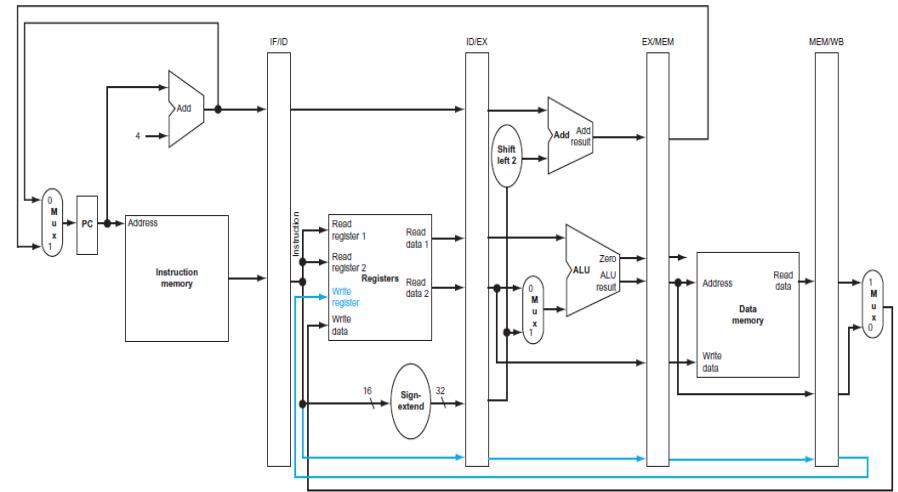


FIGURE 4.41 The corrected pipelined datapath to handle the load instruction properly. The write register number now comes from the MEM/WB pipeline register along with the data. The register number is passed from the ID pipe stage until it reaches the MEM/WB pipeline register, adding five more bits to the last three pipeline registers. This new path is shown in color.

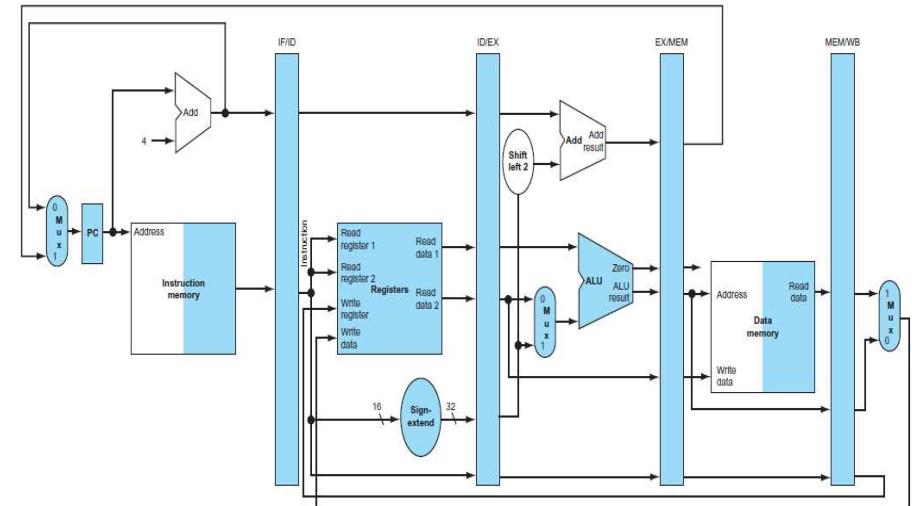


FIGURE 4.42 The portion of the datapath in Figure 4.41 that is used in all five stages of a load instruction.

Introduction to Computer Organization _ Processor Design _ Pipeline _ Datapath Hazard

Datapath with Control

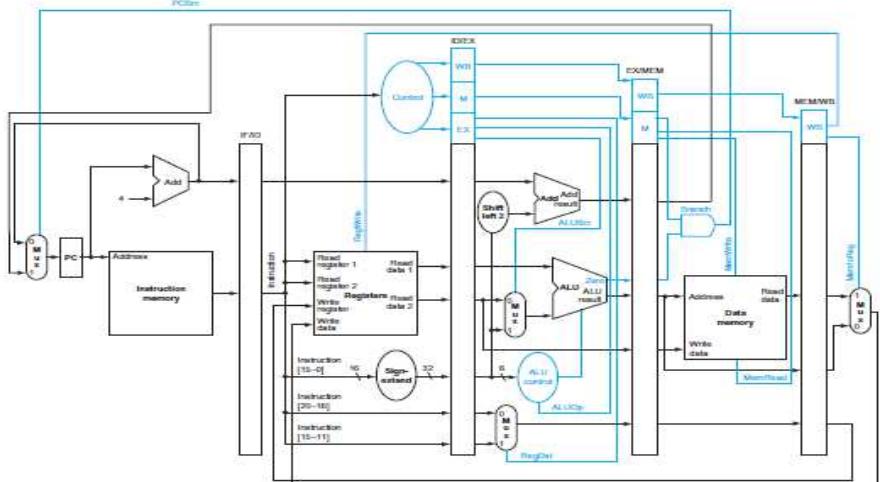


FIGURE 4.51 The pipelined datapath of Figure 4.46, with the control signals connected to the control portions of the pipeline registers. The control values for the last three stages are created during the instruction decode stage and then placed in the ID/EX pipeline register. The control lines for each pipe stage are used, and remaining control lines are then passed to the next pipeline stage.

Forwarding

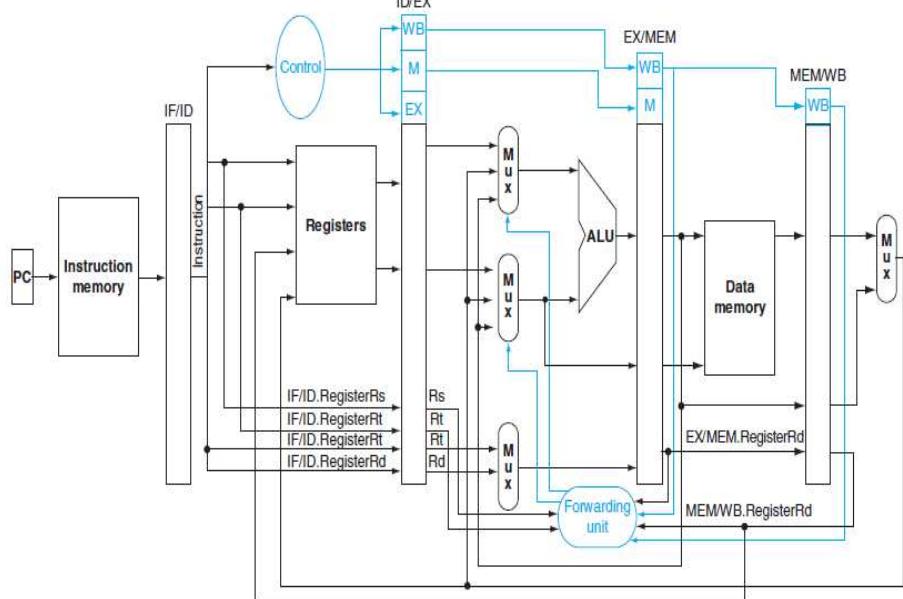


FIGURE 4.56 The datapath modified to resolve hazards via forwarding. Compared with the datapath in Figure 4.51, the additions are the multiplexors to the inputs to the ALU. This figure is a more stylized drawing, however, leaving out details from the full datapath, such as the branch hardware and the sign extension hardware.

Forwarding with Detection

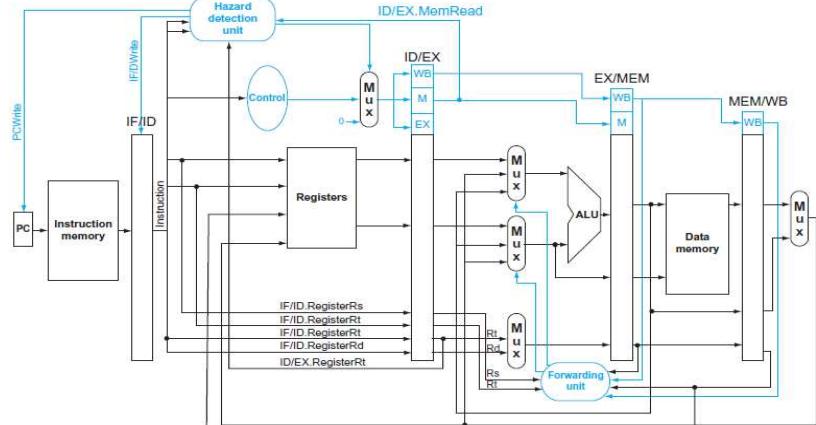
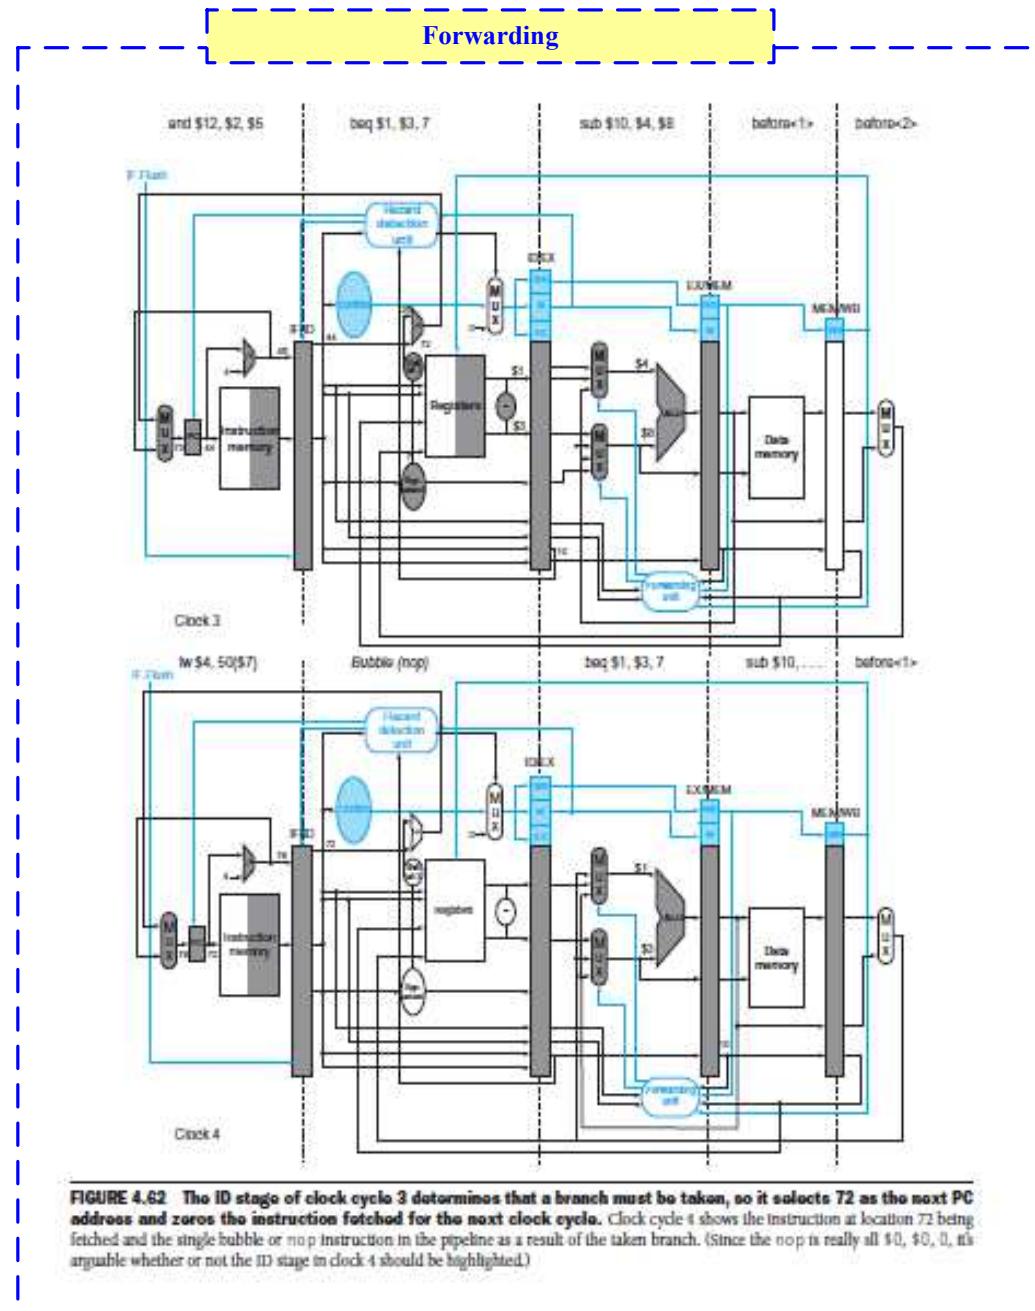
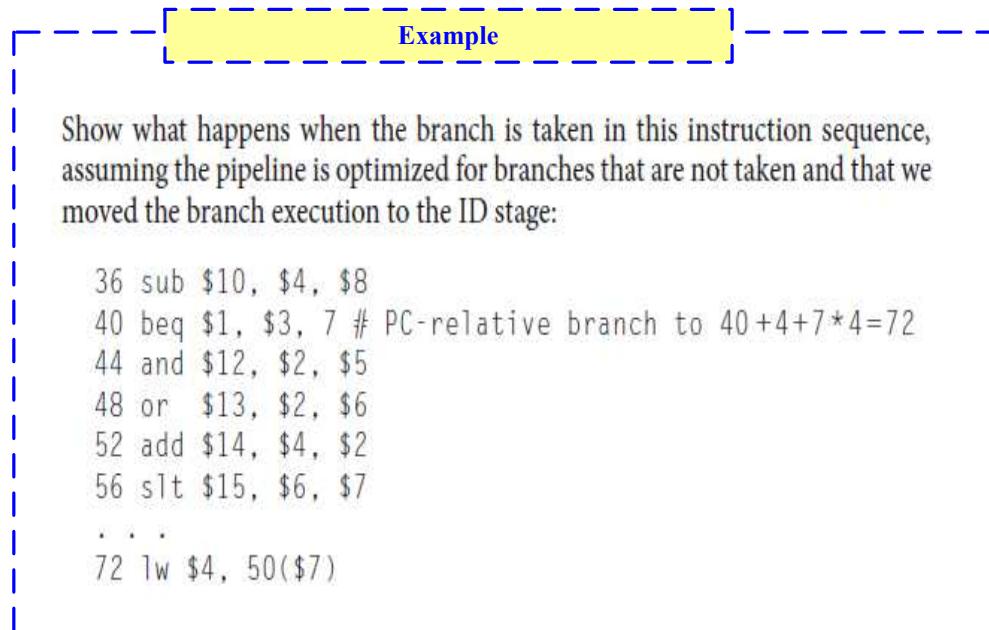
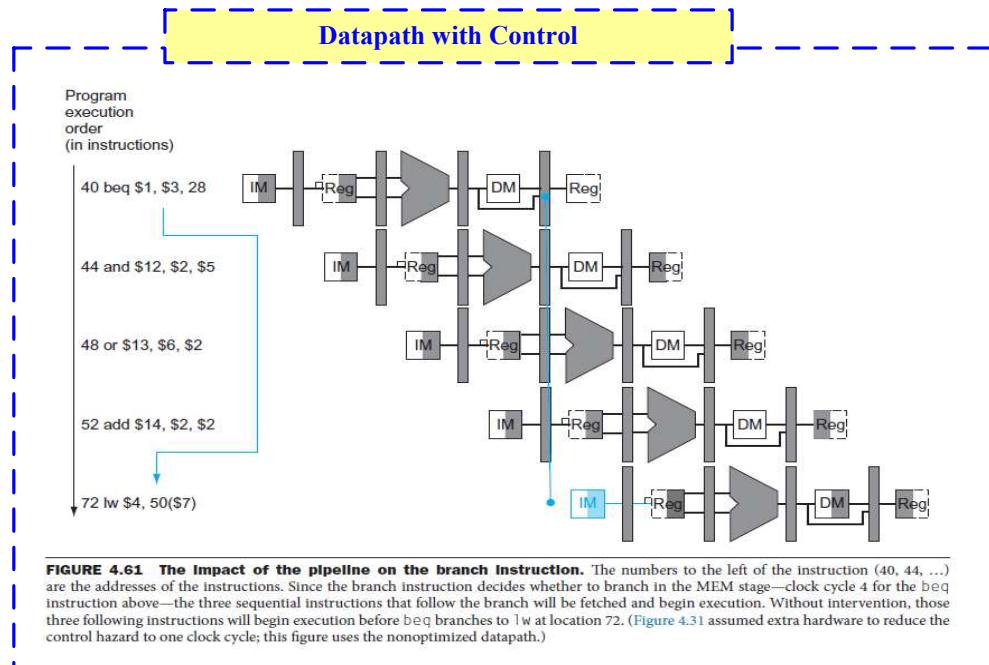


FIGURE 4.60 Pipelined control overview, showing the two multiplexors for forwarding, the hazard detection unit, and the forwarding unit. Although the ID and EX stages have been simplified—the sign-extended immediate and branch logic are missing—this drawing gives the essence of the forwarding hardware requirements.

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

FIGURE 4.55 The control values for the forwarding multiplexors in Figure 4.54. The signed immediate that is another input to the ALU is described in the *Elaboration* at the end of this section.

Introduction to Computer Organization _ Processor Design _ Pipeline _ Control Hazard



Introduction to Computer Organization _ Processor Design _ Pipeline _ Exception

Type of Event

Type of event	From where?	MIPS terminology
I/O device request	External	Interrupt
Invoke the operating system from user program	Internal	Exception
Arithmetic overflow	Internal	Exception
Using an undefined instruction	Internal	Exception
Hardware malfunctions	Either	Exception or interrupt

Exception type

Exception vector address (in hex)

Undefined instruction	8000 0000 _{hex}
Arithmetic overflow	8000 0180 _{hex}

Example

Exception in a Pipelined Computer

Given this instruction sequence,

```

40hex sub $11, $2, $4
44hex and $12, $2, $5
48hex or $13, $2, $6
4Chex add $1, $2, $1
50hex slt $15, $6, $7
54hex lw $16, 50($7)
...
    
```

Forwarding

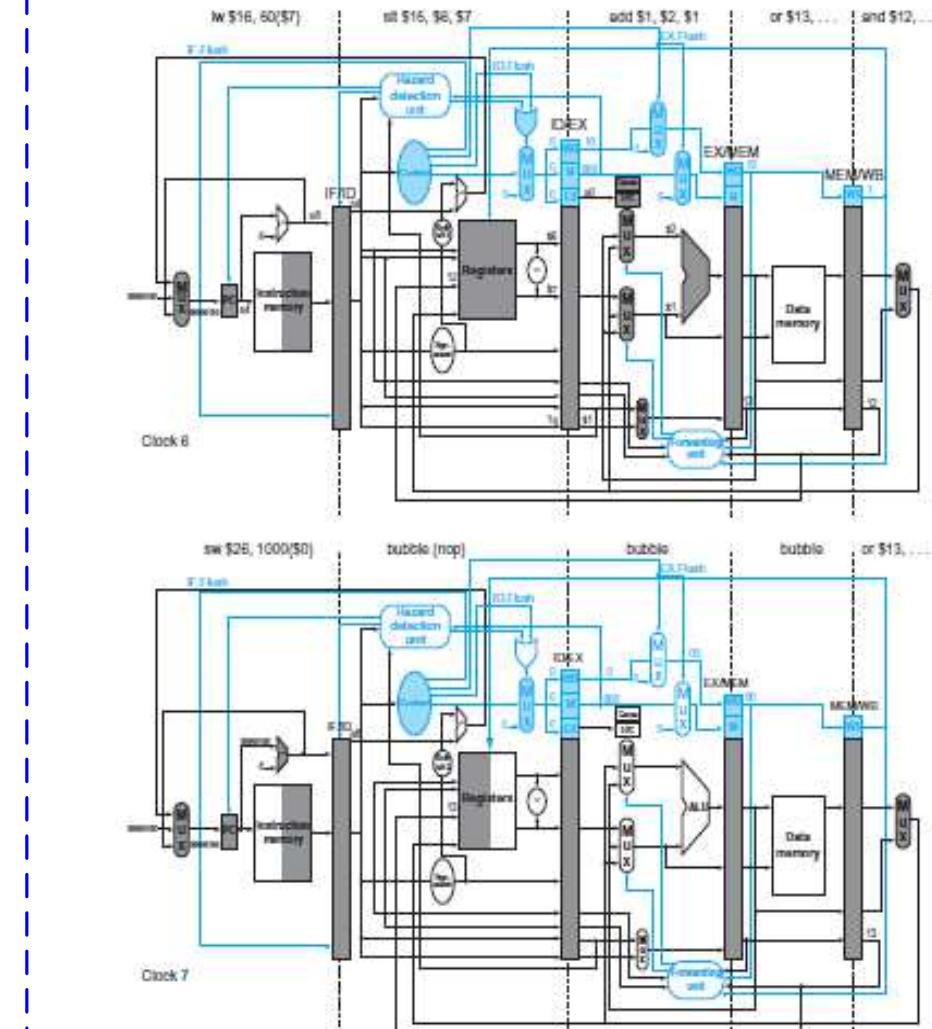


FIGURE 4.67 The result of an exception due to arithmetic overflow in the add instruction. The overflow is detected during the EX stage of clock 6, saving the address following the add in the EPC register (4C + 4 = 50_{hex}). Overflow causes all the flush signals to be set near the end of this clock cycle, deasserting control values (setting them to 0) for the add. Clock cycle 7 shows the instructions converted to bubbles in the pipeline plus the fetching of the first instruction of the exception routine—sw \$25, 1000(\$0)—from instruction location 8000 0180_{hex}. Note that the AND and OR instructions, which are prior to the add, still complete. Although not shown, the ALU overflow signal is an input to the control unit.

Introduction to Computer Organization _ Processor Design _ Enhancement

Five Classic Components of a Computer

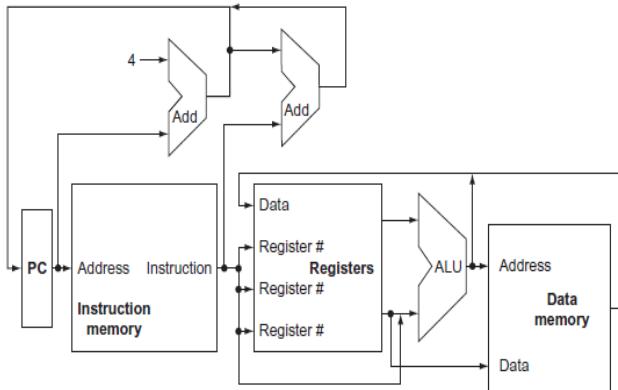


FIGURE 4.1 An abstract view of the implementation of the MIPS subset showing the major functional units and the major connections between them.

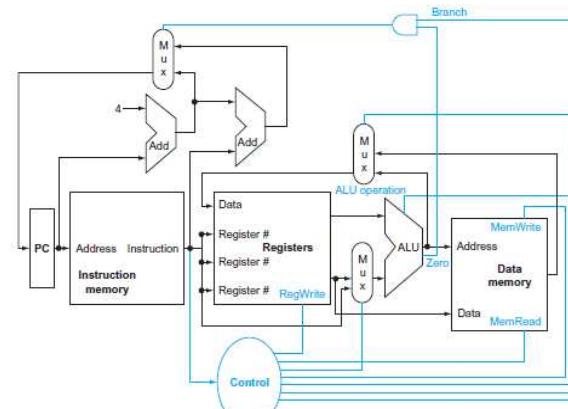


FIGURE 4.2 The basic implementation of the MIPS subset, including the necessary multiplexors and control lines.

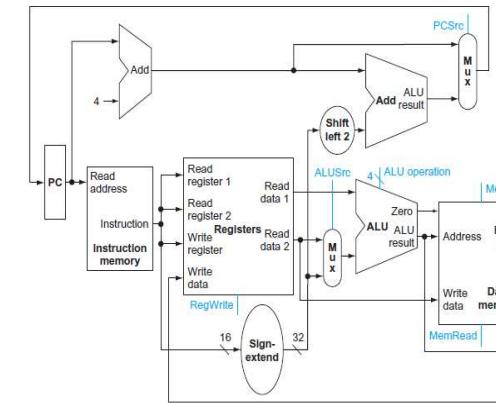


FIGURE 4.11 The simple datapath for the core MIPS architecture combines the elements required by different instruction classes.

Datapath

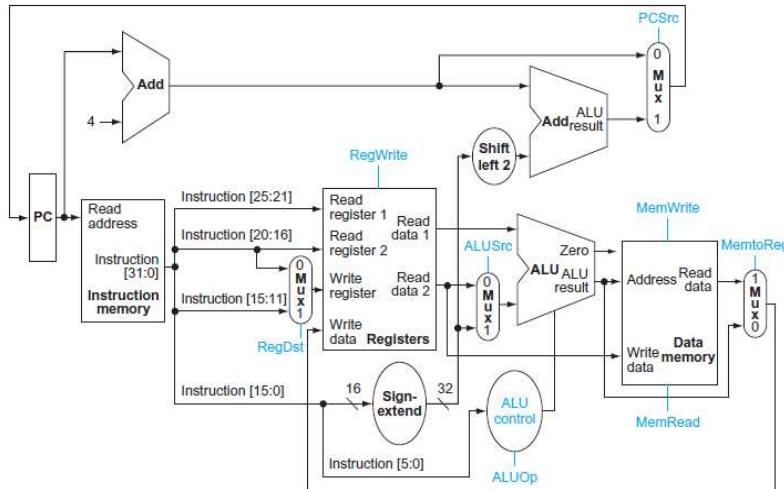


FIGURE 4.15 The datapath of Figure 4.11 with all necessary multiplexors and all control lines identified. The control lines are shown in color. The ALU control block has also been added. The PC does not require a write control, since it is written on the end of every clock cycle; the branch control logic determines whether it is written with the incremented PC or the branch target address.

Datapath with Control

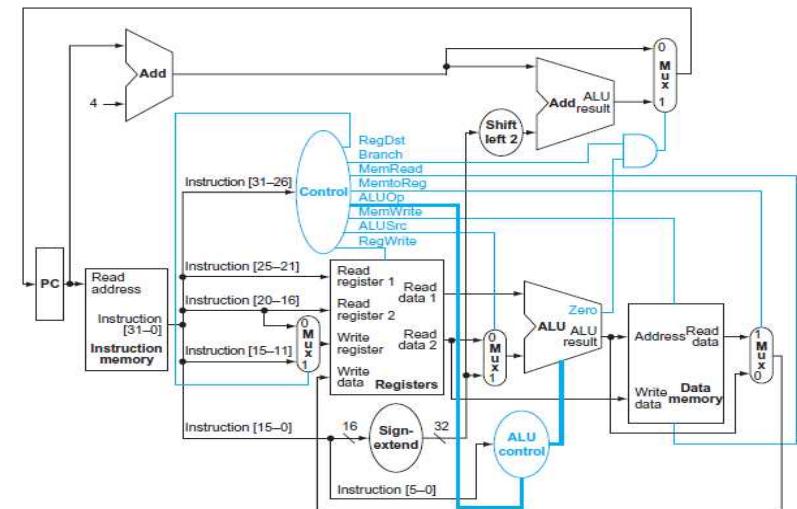


FIGURE 4.17 The simple datapath with the control unit. The input to the control unit is the 6-bit opcode field from the instruction. The outputs of the control unit consist of three 1-bit signals that are used to control multiplexors (RegDst, ALUSrc, and MemtoReg), three signals for controlling reads and writes in the register file and data memory (RegWrite, MemRead, and MemWrite), a 1-bit signal used in determining whether to possibly branch (Branch), and a 2-bit control signal for the ALU (ALUOp). An AND gate is used to combine the branch control signal and the Zero output from the ALU; the AND gate output controls the selection of the next PC. Notice that PCSrc is now a derived signal, rather than one coming directly from the control unit. Thus, we drop the signal name in subsequent figures.

Introduction to Computer Organization _ Processor Design _ R/L/B/J

R-type

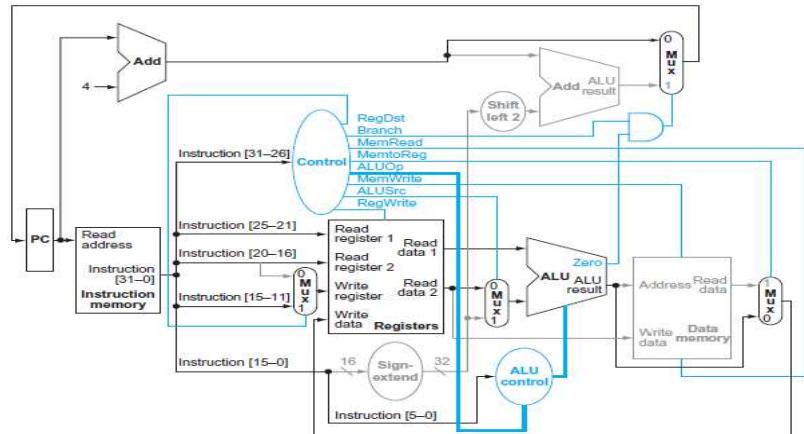


FIGURE 4.19 The datapath in operation for an R-type instruction, such as add \$t1,\$t2,\$t3. The control lines, datapath units, and connections that are active are highlighted.

Load

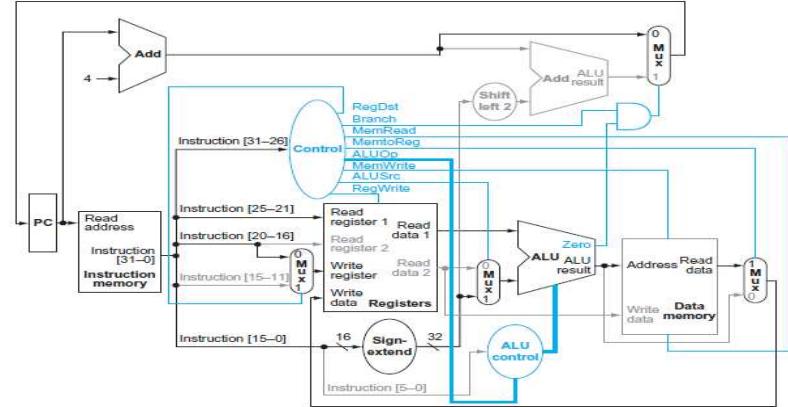


FIGURE 4.20 The datapath in operation for a load instruction. The control lines, datapath units, and connections that are active are highlighted. A store instruction would operate very similarly. The main difference would be that the memory control would indicate a write rather than a read, the second register value read would be used for the data to store, and the operation of writing the data memory value to the register file would not occur.

Jump

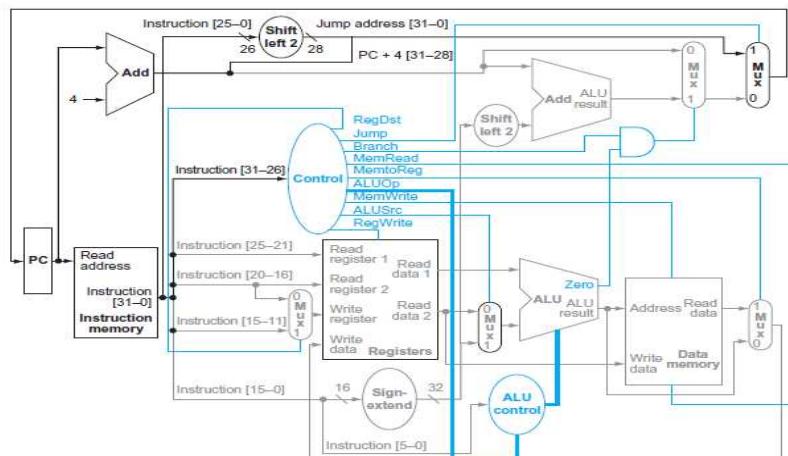


FIGURE 4.24 The simple control and datapath are extended to handle the jump instruction. An additional multiplexer (at the upper right) is used to choose between the jump target and either the branch target or the sequential instruction following this one. This multiplexer is controlled by the jump control signal. The jump target address is obtained by shifting the lower 26 bits of the jump instruction left 2 bits, effectively adding 00 as the low-order bits, and then concatenating the upper 4 bits of PC + 4 as the high-order bits, thus yielding a 32-bit address.

Branch

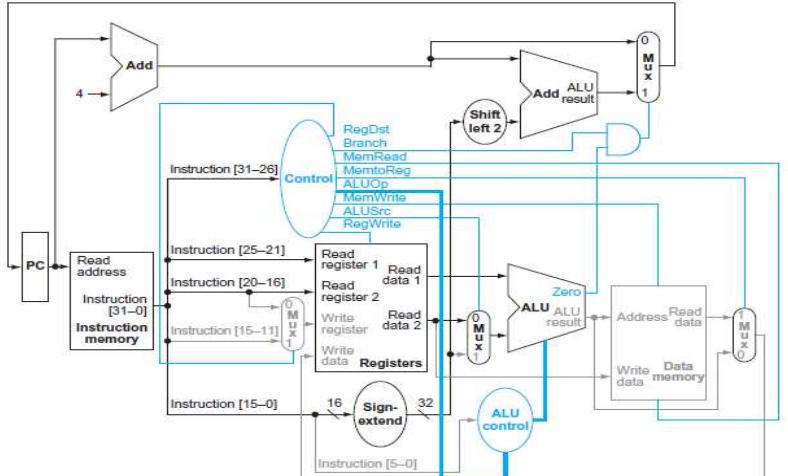


FIGURE 4.21 The datapath in operation for a branch-on-equal instruction. The control lines, datapath units, and connections that are active are highlighted. After using the register file and ALU to perform the compare, the Zero output is used to select the next program counter from between the two candidates.

Introduction to Computer Organization _ Cache Memory _Basic Concept

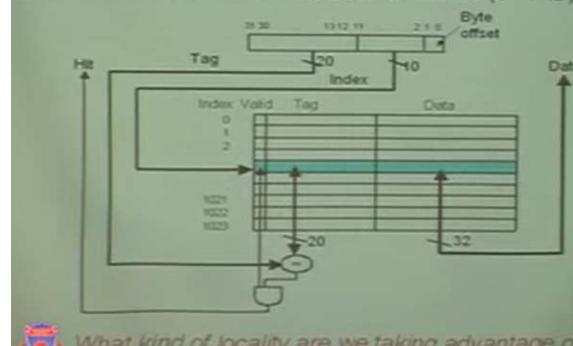
Cach Basic

Cache Basics

- Two questions to answer (in hardware):
 - Q1: How do we know if a data item is in the cache?
- Direct mapped
 - Each memory block is mapped to exactly one block in the cache
 - lots of lower level blocks must **share** blocks in the cache
 - Address mapping (to answer Q2):
(block address) modulo (# of blocks in the cache)
 - Have a **tag** associated with each cache block that contains the address information (the upper portion of the address) required to identify the block (to answer Q1)

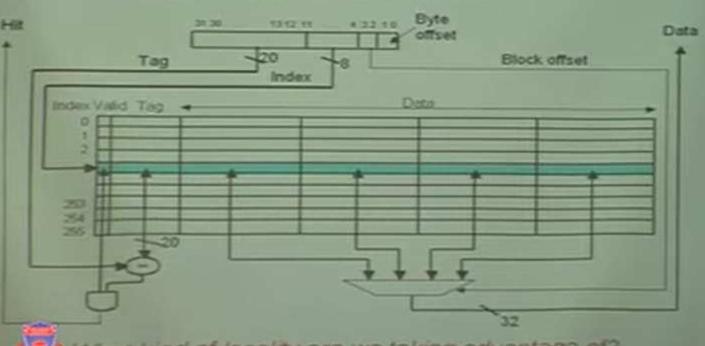
MIPS Direct Mapped Cache Example

- One word blocks, cache size = 1K words (or 4KB)



Multiword Block Direct Mapped Cache

- Four words/block, cache size = 1K words



Cache Hit and Miss

Handling Cache Hits

- Read hits (IS and DS)
 - this is what we want!
- Write hits (DS only)
 - require the cache and memory to be **consistent**
 - always write the data into both the cache block and the next level in the memory hierarchy (**write-through**)
 - writes run at the speed of the next level in the memory hierarchy – so slow! – or can use a **write buffer** and stall only if the write buffer is full
 - allow cache and memory to be **inconsistent**
 - write the data only into the cache block (write-back the cache block to the next level in the memory hierarchy when that cache block is “evicted”)
 - need a **dirty** bit for each data cache block to tell if it needs to be written back to memory when it is evicted – can use a **write buffer** to help “buffer” write-backs of dirty blocks

Sources of Cache Misses

- Compulsory (cold start or process migration, first reference):
 - First access to a block, “cold” fact of life, not a whole lot you can do about it. If you are going to run “millions” of instruction, compulsory misses are insignificant
 - Solution: increase block size (increases miss penalty; very large blocks could increase miss rate)
- Capacity:
 - Cache cannot contain all blocks accessed by the program
 - Solution: increase cache size (may increase access time)
- Conflict (collision):
 - Multiple memory locations mapped to the same cache location
 - Solution 1: increase cache size
 - Solution 2: increase associativity (stay tuned) (may increase access time)

Handling Cache Misses (Single Word Blocks)

- Read misses (IS and DS)
 - stall the pipeline, fetch the block from the next level in the memory hierarchy, install it in the cache and send the requested word to the processor, then let the pipeline resume
 - Write misses (DS only)
 - stall the pipeline, fetch the block from next level in the memory hierarchy, install it in the cache (which may involve having to evict a dirty block if using a write-back cache), write the word from the processor to the cache, then let the pipeline resume
- or
- Write allocate** – just write the word into the cache updating both the tag and data, no need to check for cache hit, no need to stall
- or
- No-write allocate** – skip the cache write (but must invalidate that cache block since it will now hold stale data) and just write the word to the write buffer (and eventually to the next memory level). 094 Spring 2014 NEED TO STALL IF THE WRITE BUFFER ISN'T FULL

Introduction to Computer Organization _ Cache Memory _ Reducing Cache Miss Rate

Improving Cache Performance

0. Reduce the time to hit in the cache

- smaller cache
- direct mapped cache
- smaller blocks
- for writes
 - no write allocate – no “hit” on cache, just write to write buffer
 - write allocate – to avoid two cycles (first check for hit, then write) pipeline writes via a delayed write buffer to cache

1. Reduce the miss rate

- bigger cache
- more flexible placement (increase associativity)
- larger blocks (16 to 64 bytes typical)
- victim cache – small buffer holding most recently discarded blocks

2. Reduce the miss penalty

- smaller blocks
- use a write buffer to hold dirty blocks being replaced so don't have to wait for the write to complete before reading
- check write buffer (and/or victim cache) on read miss – may get lucky
- for large blocks fetch critical word first
- use multiple cache levels – L2 cache not tied to CPU clock rate
- faster backing store/improved memory bandwidth
 - wider buses
 - memory interleaving, DDR SDRAMs

Summary: The Cache Design Space

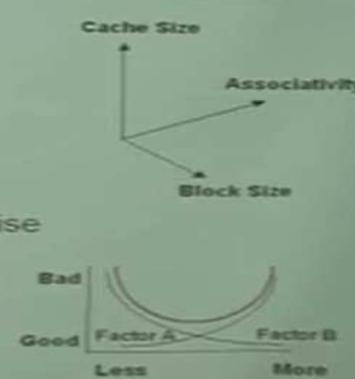
□ Several interacting dimensions

- cache size
- block size
- associativity
- replacement policy
- write-through vs write-back
- write allocation

□ The optimal choice is a compromise

- depends on access characteristics
 - workload
 - use (I-cache, D-cache, TLB)
- depends on technology / cost

• Simplicity often wins



Reducing Cache Miss Rate

Reducing Cache Miss Rates #1

1. Allow more flexible block placement
 - In a direct mapped cache a memory block maps to exactly one cache block
 - At the other extreme, could allow a memory block to be mapped to any cache block – fully associative cache
 - A compromise is to divide the cache into sets each of which consists of n “ways” (n-way set associative). A memory block maps to a unique set (specified by the index field) and can be placed in any way of that set (so there are n choices)

(block address) modulo (# sets in the cache)

Reducing Cache Miss Rates #2

1. Use multiple levels of caches

- With advancing technology have more than enough room on the die for bigger L1 caches or for a second level of caches – normally a unified L2 cache (i.e., it holds both instructions and data) and in some cases even a unified L3 cache
- For our example, CPI_{ideal} of 2, 100 cycle miss penalty (to main memory) and a 25 cycle miss penalty (to UL2S), 36% load/stores, a 2% (4%) L1 IS (DS) miss rate, add a 0.5% UL2S miss rate

$$\begin{aligned} \text{CPI}_{\text{stalls}} = & 2 + .02 \times 25 + .36 \times .04 \times 25 + .005 \times 100 + \\ & .36 \times .005 \times 100 = 3.54 \end{aligned}$$

(as compared to 5.44 with no L2\$)

Introduction to Computer Organization _ TLB _ Basic Concept

Translation Lookaside Buffers

Translation Lookaside Buffers (TLBs)

- Just like any other cache, the TLB can be organized as fully associative, set associative, or direct mapped

V	Virtual Page #	Physical Page #	Dirty	Ref	Access

- TLB access time is typically smaller than cache access time (because TLBs are much smaller than caches)
 - TLBs are typically not more than 512 entries even on high-end machines

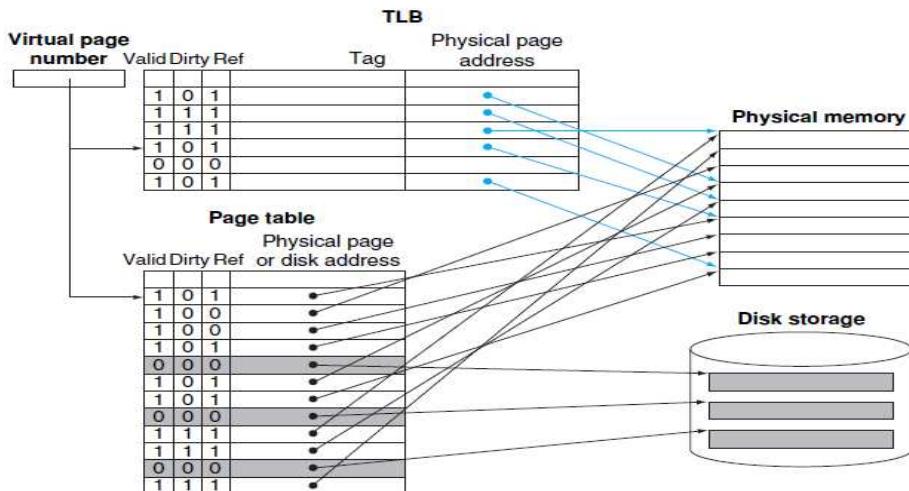
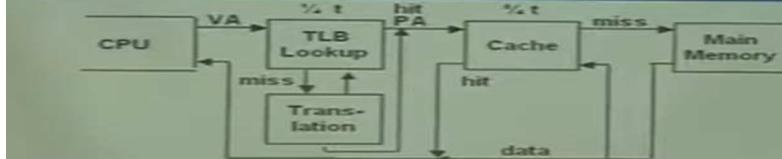


FIGURE 5.29 The TLB acts as a cache of the page table for the entries that map to physical pages only. The TLB contains a subset of the virtual-to-physical page mappings that are in the page table. The TLB mappings are shown in color. Because the TLB is a cache, it must have a tag field. If there is no matching entry in the TLB for a page, the page table must be examined. The page table either supplies a physical page number for the page (which can then be used to build a TLB entry) or indicates that the page resides on disk, in which case a page fault occurs. Since the page table has an entry for every virtual page, no tag field is needed; in other words, unlike a TLB, a page table is not a cache.

TLB in Memory Hierarchy

A TLB in the Memory Hierarchy



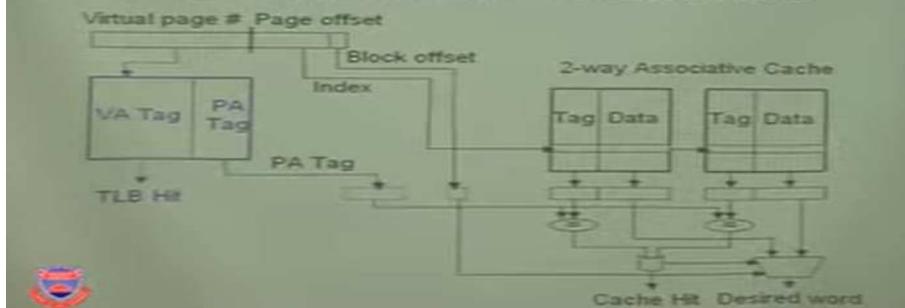
- A TLB miss – is it a page fault or merely a TLB miss?
 - If the page is loaded into main memory, then the TLB miss can be handled (in hardware or software) by loading the translation information from the page table into the TLB
 - Takes 10's of cycles to find and load the translation info into the TLB
 - If the page is not in main memory, then it's a true page fault
 - Takes 1.000.000's of cycles to service a page fault
- TLB misses are much more frequent than true page faults.

TLB Event Combinations

TLB	Page Table	Cache	Possible? Under what circumstances?
Hit	Hit	Hit	Yes – what we want!
Hit	Hit	Miss	Yes – although the page table is not checked if the TLB hits
Miss	Hit	Hit	Yes – TLB miss, PA in page table
Miss	Hit	Miss	Yes – TLB miss, PA in page table, but data not in cache
Miss	Miss	Miss	Yes – page fault
Hit	Miss	Miss/Hit	Impossible – TLB-translation-not-possible if page is not present in memory
Miss	Miss	Hit	Impossible – data-not-allowed-in-cache-if-page is not in memory

Reducing Translation Time

- Can overlap the cache access with the TLB access
 - Works when the high order bits of the VA are used to access the TLB while the low order bits are used as index into cache



Introduction to Computer Organization _ TLB _ SW&HW Boundary

Boundary of S/W and H/w

- What parts of the virtual to physical address translation is done by or assisted by the hardware?
 - Translation Lookaside Buffer (TLB) that caches the recent translations
 - TLB access time is part of the cache hit time
 - May allow an extra stage in the pipeline for TLB access
 - Page table storage, fault detection and updating
 - Page faults result in interrupts (precise) that are then handled by the OS
 - Hardware must support (i.e., update appropriately) Dirty and Reference bits (e.g., ~LRU) in the Page Tables
 - Disk placement
 - Bootstrap (e.g., out of disk sector 0) so the system can service a limited number of page faults before the OS is even loaded

Summary

The Principle of Locality:

- Program likely to access a relatively small portion of the address space at any instant of time.
 - Temporal Locality: Locality in Time
 - Spatial Locality: Locality in Space

Caches, TLBs, Virtual Memory all understood by examining how they deal with the four questions

- Where can entry be placed?
- How is entry found?
- What entry is replaced on miss?
- How are writes handled?

Page tables map virtual address to physical address

TLBs are important for fast translation

Issues

Q1&Q2: Where can a entry be placed/found?

	# of sets	Entries per set
Direct mapped	# of entries	1
Set associative	(# of entries)/ associativity	Associativity (typically 2 to 16)
Fully associative	1	# of entries

	Location method	# of comparisons
Direct mapped	Index	1
Set associative	Index the set; compare set's tags	Degree of associativity
Fully associative	Compare all entries' tags Separate lookup (page) table	# of entries

Q3: Which entry should be replaced on a miss?

- Easy for direct mapped – only one choice
- Set associative or fully associative
 - Random
 - LRU (Least Recently Used)
- For a 2-way set associative, random replacement has a miss rate about 1.1 times higher than LRU
- LRU is too costly to implement for high levels of associativity (> 4-way) since tracking the usage information is costly

Q4: What happens on a write?

Write-through – The information is written to the entry in the current memory level and to the entry in the next level of the memory hierarchy

- Always combined with a write buffer so write waits to next level in memory can be eliminated (as long as the write buffer doesn't fill)

Write-back – The information is written only to the entry in the current memory level. The modified entry is written to next level of memory only when it is replaced.

- Need a dirty bit to keep track of whether the entry is clean or dirty
- Virtual memory systems always use write-back of dirty pages to disk

Pros and cons of each?

- Write-through: read misses don't result in writes (so are simpler and cheaper), easier to implement

Write-back: writes run at the speed of the cache; repeated writes require only one write to lower level

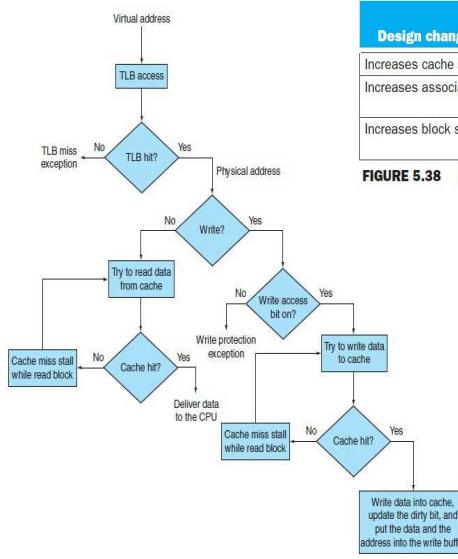


FIGURE 5.31 Processing a read or a write-through in the Intrinsity FastMATH TLB and cache.

Design change	Effect on miss rate	Possible negative performance effect
Increases cache size	Decreases capacity misses	May increase access time
Increases associativity	Decreases miss rate due to conflict misses	May increase access time
Increases block size	Decreases miss rate for a wide range of block sizes due to spatial locality	Increases miss penalty. Very large block could increase miss rate

FIGURE 5.38 Memory hierarchy design challenges.

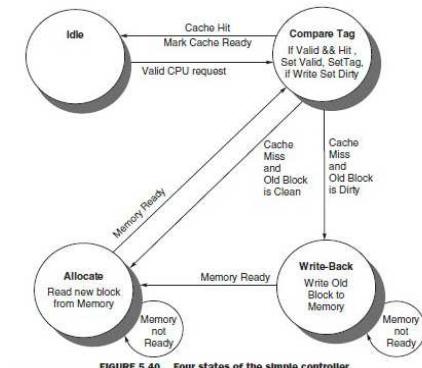


FIGURE 5.40 Four states of the simple controller.

Introduction to Computer Organization _ Storage Management _ NAS and SAN

Storage-device Hierarchy

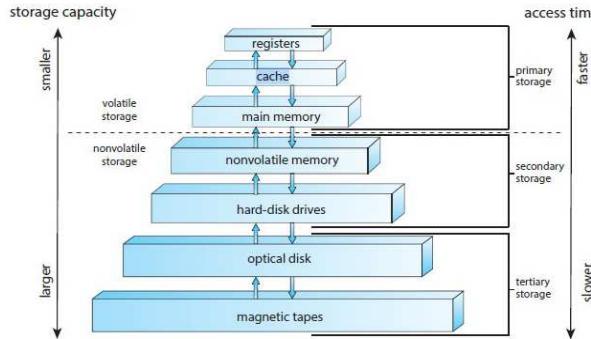
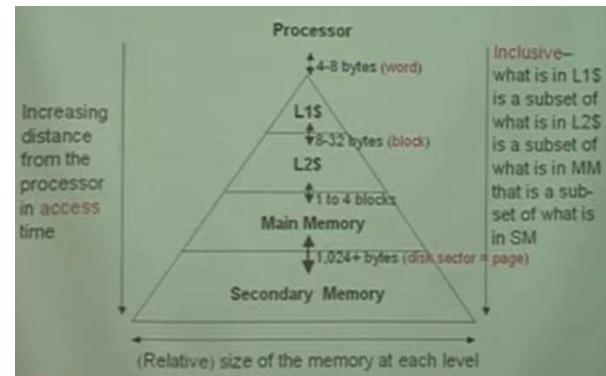
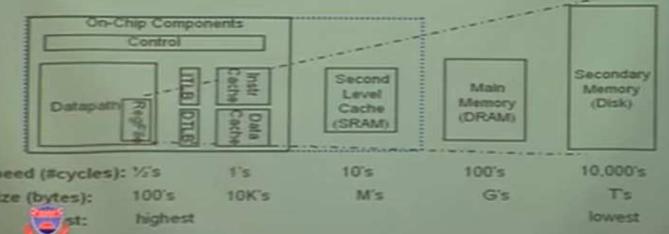


Figure 1.6 Storage-device hierarchy.



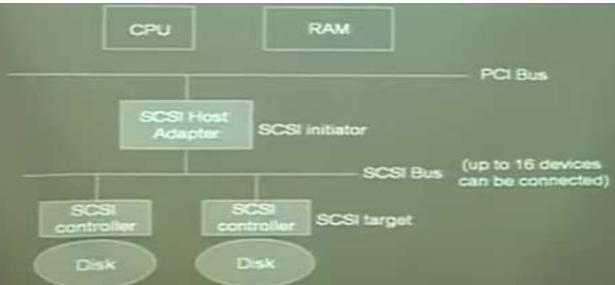
A Typical Memory Hierarchy

- Take advantage of the principle of locality to present the user with as much memory as is available in the cheapest technology at the speed offered by the fastest technology



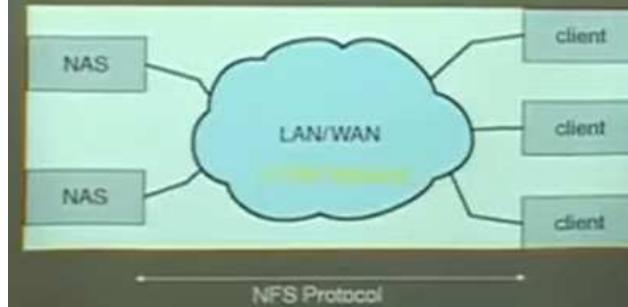
Storage Attachement _ NAS vs SAN

- Host-attached storage accessed through I/O ports talking to I/O busses
- Attachment technologies and protocols
 - IDE
 - SCSI
 - Fiber Channel
- SCSI itself is a bus, up to 16 devices on one cable, SCSI initiator requests operation and SCSI targets perform tasks
 - Each target can have up to 8 logical units (disks attached to device controller)
- FC (fiber channel) is high-speed serial architecture
 - Can be switched fabric with 24-bit address space – the basis of storage area networks (SANs) in which many hosts attach to many storage units
 - Can be arbitrated loop (FC-AL) of 126 devices



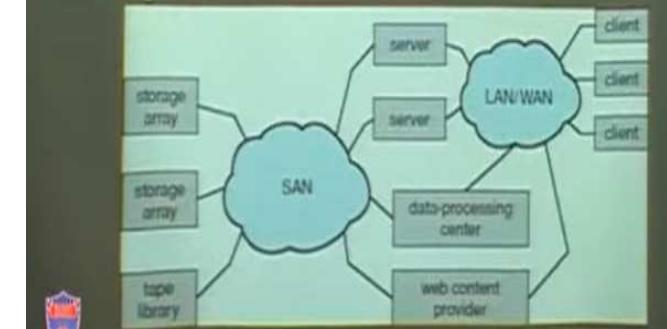
Network Attached Storage

- Network-attached storage (NAS) is storage made available over a network rather than over a local connection (such as a bus)
- NFS and CIFS are common protocols used for network attached storage
 - We use those protocols to access remote storage that is connected to a network.
- Implemented via remote procedure calls (RPCs) between host and storage
- New iSCSI protocol uses IP network to carry the SCSI protocol
 - SCSI – SCSI bus/cable (local/host attached)
 - iSCSI – TCP/IP network (network attached)



Storage Area Network

- Common in large storage environments (and becoming more common)
- Multiple hosts attached to multiple storage arrays – flexible
- Uses a different communication infrastructure (SAN) than the common networking infrastructure



Introduction to Computer Organization _ Storage Management _ Disk Management

Magnetic Disk / Hard Disk

Overview of Mass Storage Systems: Magnetic Disks

- Magnetic disks provide bulk of secondary storage of modern computers
 - Drives rotate at 60 to 200 times per second
 - Transfer rate** is rate at which data flow between drive and computer
 - Positioning time (random-access time)** is time to move disk arm to desired cylinder (**seek time**) and time for desired sector to rotate under the disk head (**rotational latency**)
 - Head crash** results from disk head making contact with the disk surface
 - That's bad
- Disks can be removable
- Drive attached to computer via I/O bus
 - Busses vary, including EIDE, ATA, SATA, USB, Fibre Channel, SCSI
 - Host controller** in computer uses bus to talk to **disk controller** built into drive or storage array

Magnetic Disk Characteristic

- Disk read/write components
 - Seek time:** position the head over the proper track (3 to 13 ms avg)
 - due to locality of disk references the actual average seek time may be only 25% to 33% of the advertised number
 - Rotational latency:** wait for the desired sector to rotate under the head (1/2 of 1 rotation/RPM, converted to ms)
 - $0.5/5400\text{RPM} = 5.6 \text{ ms}$ to $0.5/15000\text{RPM} = 2.0 \text{ ms}$
 - Transfer time:** transfer a block of bits (one or more sectors) under the head to the disk controller's cache (70 to 125 MB/s are typical disk transfer rates in 2008)
 - disk controller's cache takes advantage of spatial locality in disk accesses. Cache transfer rates are much faster (e.g. 375 MB/s)
- Controller time: the overhead the disk controller imposes in performing a disk I/O access (typically <0.2 ms)

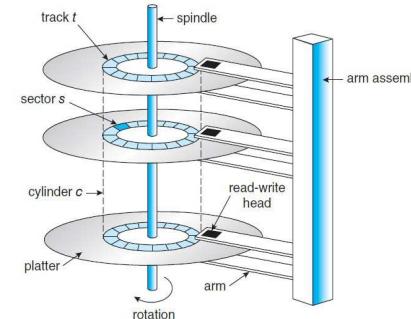


Figure 11.1 HDD moving-head disk mechanism.

Disk Structure

- Disk drives are addressed as large 1-dimensional arrays of logical blocks, where the logical block is the smallest unit of transfer.
- The 1-dimensional array of logical blocks is mapped into the sectors of the disk sequentially.
 - Sector 0 is the first sector of the first track on the outermost cylinder.
 - Mapping proceeds in order through that track, then the rest of the tracks in that cylinder, and then through the rest of the cylinders from outermost to innermost.

Disk Scheduling

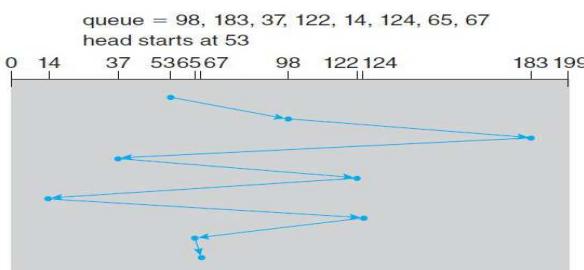


Figure 11.6 FCFS disk scheduling.

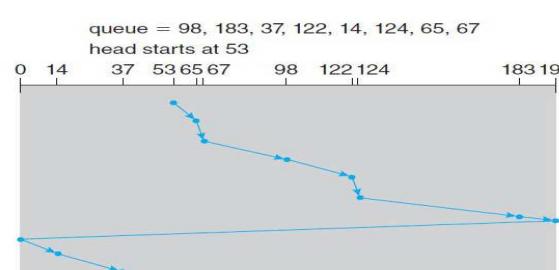


Figure 11.8 C-SCAN disk scheduling.

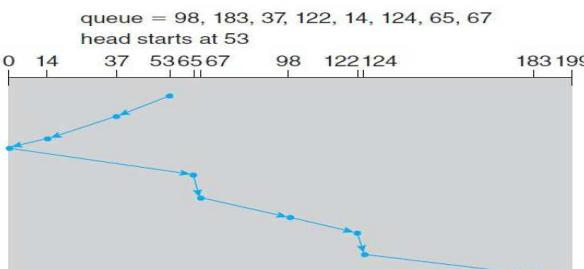
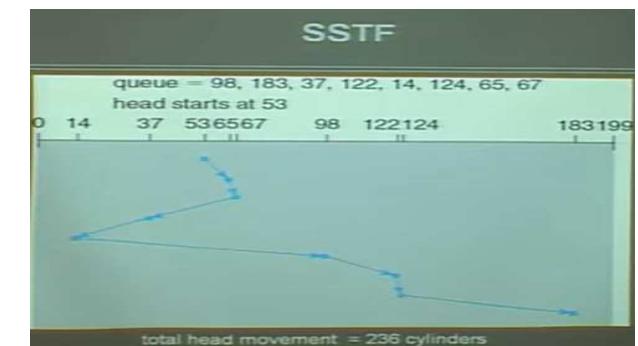
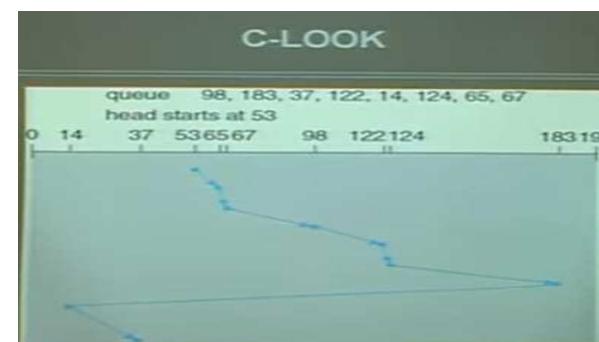


Figure 11.7 SCAN disk scheduling.



Selecting a Disk-Scheduling Algorithm

- SSTF is common and has a natural appeal
- SCAN and C-SCAN perform better for systems that place a heavy load on the disk.
- Performance depends on the number and types of requests.
- Requests for disk service can be influenced by the file-allocation method.
- The disk-scheduling algorithm should be written as a separate module of the operating system, allowing it to be replaced with a different algorithm if necessary.
- Either SSTF or LOOK is a reasonable choice for the default algorithm.

Introduction to Computer Organization _ Storage Management _ DAID

Mass Storage

Dependability, Reliability, Availability

- Reliability – measured by the mean time to failure (MTTF). Service interruption is measured by mean time to repair (MTTR)
- Availability – a measure of service accomplishment
 $\text{Availability} = \frac{\text{MTTF}}{\text{MTTF} + \text{MTTR}}$
- To increase MTTF, either improve the quality of the components or design the system to continue operating in the presence of faulty components
 - Fault avoidance: preventing fault occurrence by construction
 - Fault tolerance: using redundancy to correct or bypass faulty components (hardware)
 - Fault detection versus fault correction
 - Permanent faults versus transient faults

RAIDs: Disk Arrays

Redundant Array of Inexpensive Disks



Arrays of small and inexpensive disks

- Increase potential throughput by having many disk drives
 - Data is spread over multiple disk
 - Multiple accesses are made to several disks at a time

Reliability is lower than a single disk

But availability can be improved by adding redundant disks (RAID)

Lost information can be reconstructed from redundant information
 MTTR: mean time to repair is in the order of hours
 MTTF: mean time to failure of disks is tens of years

Summary

- Four components of disk access time:
 - Seek Time: advertised to be 3 to 14 ms but lower in real systems
 - Rotational Latency: 5.6 ms at 5400 RPM and 2.0 ms at 15000 RPM
 - Transfer Time: 30 to 80 MB/s
 - Controller Time: typically less than .2 ms
- RAIDS can be used to improve availability
 - RAID 1 and RAID 5 – widely used in servers, one estimate is that 80% of disks in servers are RAIDS
 - RAID 0+1 (mirroring) – EMC, HP/Tandem, IBM
 - RAID 3 – Storage Concepts
 - RAID 4 – Network Appliance
- RAIDS have enough redundancy to allow continuous operation, but hot swapping (replacement while system is running) is challenging

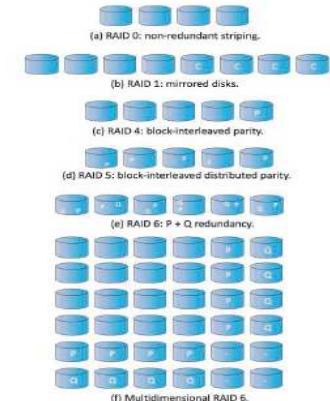
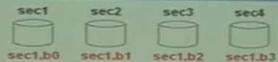


Figure 11.15 RAID levels.

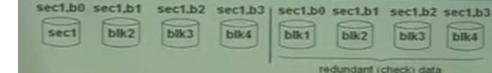
RAID Level

RAID: Level 0 (No Redundancy, but Striping)



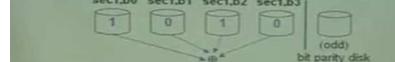
- Multiple smaller disks as opposed to one big disk
 - Spreading the sector over multiple disks – **striping** – means that multiple blocks can be accessed in parallel increasing the performance (i.e. throughput)
 - this 4 disk system gives four times the throughput of a 1 disk system
 - Same cost as one big disk – assuming 4 small disks cost the same as one big disk
- No redundancy, so what if one disk fails?
 - Failure of one or more disks is more likely as the number of disks in the system increases

RAID: Level 0+1 (Striping with Mirroring)



- Combines the best of RAID 0 and RAID 1, data is striped across four disks and mirrored to four disks
 - Four times the throughput (due to striping)
 - # redundant disks = # of data disks, so reliability doubles the cost
 - writes have to be made to both sets of disks, so writes would be only 1/2 the performance of a RAID 0 with 8 disks
- What if one disk fails?
 - If a disk fails, the system just goes to the "mirror" for the data

RAID: Level 3 (Bit-Interleaved Parity)



- Cost of higher availability is reduced to 1/N where N is the number of disks in a protection group
 - # redundant disks = 1 × # of protection groups
 - Reads and writes must access all disks
 - writing new data to the data disk as well as computing and writing the parity disk, means reading the other disks, so that the parity disk can be updated
- Can tolerate limited (single) disk failure, since the data can be reconstructed
 - Reads require reading all the operational data disks as well as the parity disk to calculate the missing data stored on the failed disk

RAID: Level 4 (Block-Interleaved Parity)

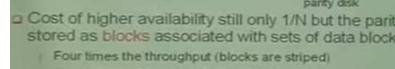


- Cost of higher availability still only 1/N but the parity is stored as blocks associated with sets of data blocks
 - Four times the throughput (blocks are striped)
 - # redundant disks = 1 × # of protection groups
 - Supports "small reads" and "small writes" (reads and writes that go to just one (or a few) data disk in a protection group, not to all)
 - by watching which bits change when writing new information, need only to change the corresponding bits on the parity disk (read-modify-write)
 - the parity disk must be updated on every write, so it is a bottleneck for back-to-back writes
- Tolerate limited disk failure, since the data can be reconstructed

Small Writes



- RAID 3 writes
 - New D1 data
 - 3 reads and 2 writes involving all the disks



- RAID 4 small writes
 - New D1 data
 - 2 reads and 2 writes involving just two disks

RAID: Level 5 (Distributed Block-Interleaved Parity)



- Cost of higher availability still only 1/N but the parity block can be located on any of the disks so there is no single bottleneck for writes
 - Still four times the throughput (block striping)
 - # redundant disks = 1 × # of protection groups
 - Supports "small reads" and "small writes" (reads and writes that go to just one (or a few) data disk in a protection group, not to all)
 - Allows multiple simultaneous writes as long as the accompanying parity blocks are not located on the same disk
- Tolerate limited disk failure, since the data can be reconstructed

Introduction to Computer Organization _ I/O System _ Basic Concept

I/O System

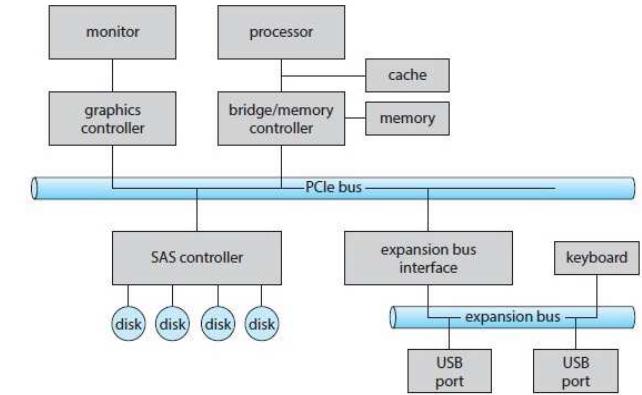
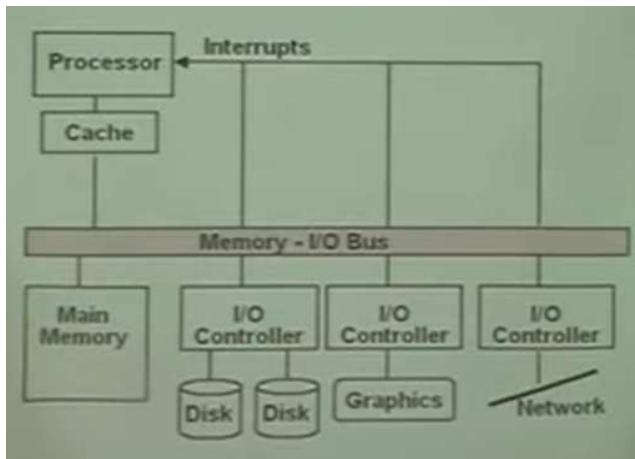
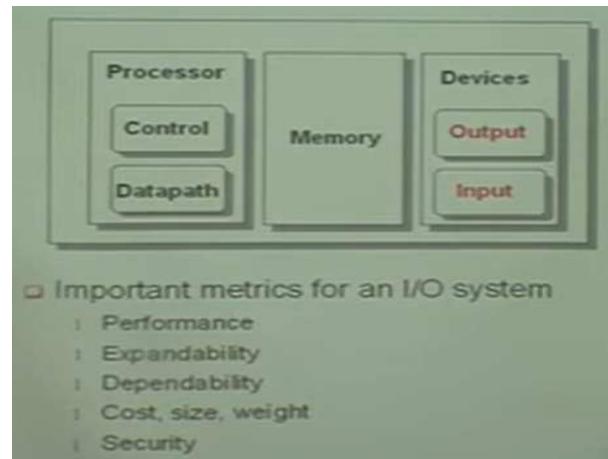
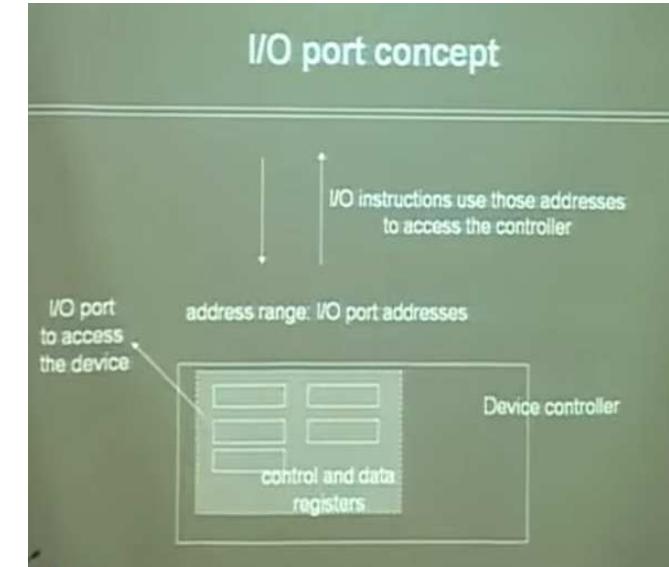
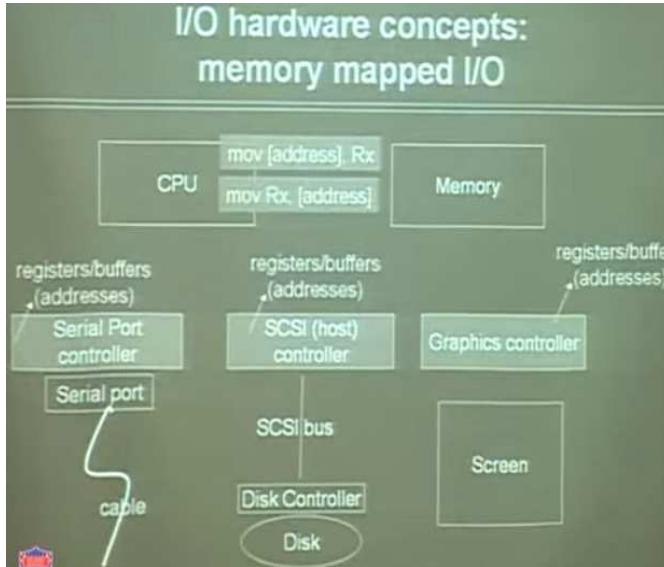
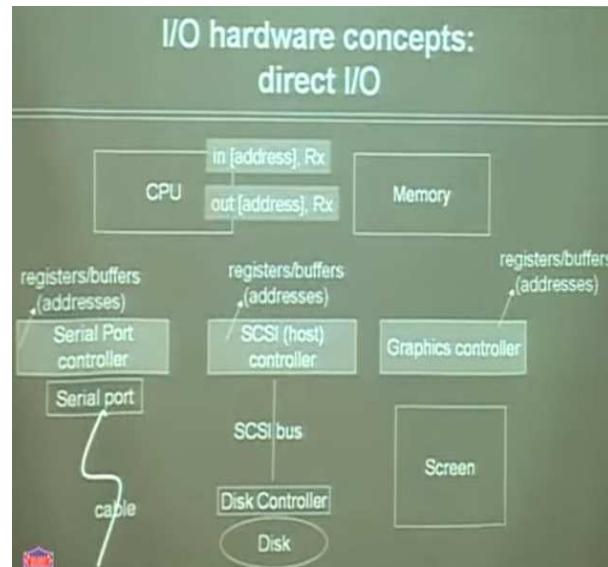


Figure 12.1 A typical PC bus structure.

I/O Concept



Introduction to Computer Organization _ I/O System _ Performance

I/O Performance

Input and Output Devices

- I/O devices are incredibly diverse with respect to
 - Behavior – input, output or storage
 - Partner – human or machine
 - Data rate – the peak rate at which data can be transferred between the I/O device and the main memory or processor

Device	Behavior	Partner	Data rate (Mb/s)
Keyboard	input	human	0.0001
Mouse	input	human	0.0038
Laser printer	output	human	3.2000
Magnetic disk	storage	machine	800.0000-3000.0000
Graphics display	output	human	800.0000-8000.0000
Network/LAN	input or output	machine	100.0000-10000.0000

↑ rates of transfer

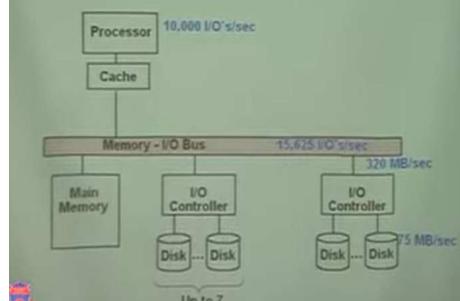
I/O Performance Measures

- I/O bandwidth (throughput) – amount of information that can be input (output) and communicated across an interconnect (e.g., a bus) to the processor/memory (I/O device) per unit time
 - How much data can we move through the system in a certain time?
 - How many I/O operations can we do per unit time?
- I/O response time (latency) – the total elapsed time to accomplish an input or output operation
 - An especially important performance metric in real-time systems
- Many applications require both high throughput and short response times

I/O System Performance

- Designing an I/O system to meet a set of bandwidth and/or latency constraints means
 - Finding the weakest link in the I/O system – the component that constrains the design
 - The processor and memory system ?
 - The underlying interconnection (i.e., bus) ?
 - The I/O controllers ?
 - The I/O devices themselves ?
 - (Re)configuring the weakest link to meet the bandwidth and/or latency requirements
 - Determining requirements for the rest of the components and (re)configuring them to support this latency and/or bandwidth

Disk I/O System Example



I/O System Performance Example

- A disk workload consisting of 64KB reads and writes where the user program executes 200,000 instructions per disk I/O operation and a processor that sustains 3 billion instr/sec and averages 100,000 OS instructions to handle a disk I/O operation

The maximum disk I/O rate (# I/O/sec) of the processor is

$$\frac{\text{Instr execution rate}}{\text{Instr per I/O}} = \frac{3 \times 10^9}{(200 + 100) \times 10^3} \approx 10,000 \text{ I/O/sec}$$

a memory-I/O bus that sustains a transfer rate of 1000 MB/s

$$\text{Each disk I/O reads/writes 64 KB so the maximum I/O rate of the bus is } \frac{\text{Bus bandwidth}}{\text{Bytes per I/O}} = \frac{1000 \times 10^6}{64 \times 10^3} \approx 15.625 \text{ I/O/sec}$$

SCSI disk I/O controllers with a DMA transfer rate of 320 MB/sec that can accommodate up to 7 disks per controller

disk drives with a read/write bandwidth of 75 MB/sec and an average seek plus rotational latency of 6 ms

What is the maximum sustainable I/O rate and what is the number of disks and SCSI controllers required to achieve that rate?

I/O System Performance Example, Con't

So the processor is the bottleneck, not the Memory-I/O bus

disk drives with a read/write bandwidth of 75 MB/sec and an average seek plus rotational latency of 6 ms

$$\text{Disk I/O read/write time} = \text{seek} + \text{rotational time} + \text{transfer time} = 6\text{ms} + 64\text{KB}/(75\text{MB/sec}) = 6.9\text{ms}$$

Thus each disk can complete 1000ms/6.9ms or 146 I/O's per second. To satisfy the processor requires 10,000 I/O's per second.

$$10,000/146 = 69 \text{ disks}$$

To calculate the number of SCSI disk controllers, we need to know the average transfer rate per disk, in order to ensure we can put the maximum of 7 disks per SCSI controller and that a disk controller won't saturate the memory-I/O bus during a DMA transfer

$$\text{Disk transfer rate} = (\text{transfer size} / \text{transfer time}) = 64\text{KB}/6.9\text{ms} = 9.56 \text{ MB/sec}$$

Thus 7 disks won't saturate either the SCSI controller (with a maximum transfer rate of 320 MB/sec) or the memory-I/O bus (1000 MB/sec). This means we will need 6/7 or 10 SCSI controllers

I/O Interface

I/O Transactions

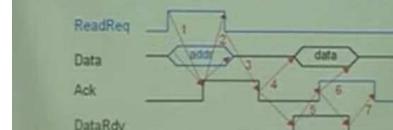
- An I/O transaction is a sequence of operations over the interconnect that includes a request and may include a response either of which may carry data. A transaction is initiated by a single request and may take many individual bus operations. An I/O transaction typically includes two parts
 - Sending the address
 - Receiving or sending the data

Bus transactions are defined by what they do to memory

- A **read** transaction reads data from memory (to either the processor or an I/O device)
- A **write** transaction writes data to the memory (from either the processor or an I/O device)

Asynchronous Bus Handshaking Protocol

- Output (read) data from memory to an I/O device



I/O device signals a request by raising ReadReq and putting the addr on the data lines

Memory sees ReadReq, reads addr from data lines, and raises Ack

I/O device sees Ack and releases the ReadReq and data lines

Memory sees ReadReq go low and drops Ack

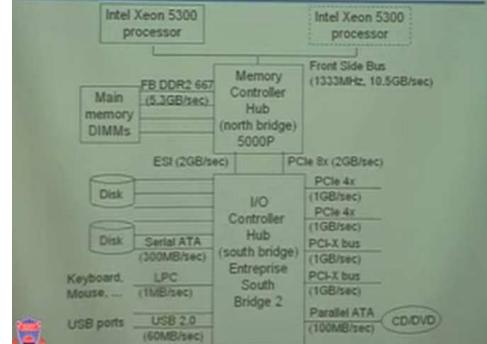
When memory has data ready, it places it on data lines and raises DataRdy

I/O device sees DataRdy, reads the data from data lines, and raises Ack

Memory sees Ack, releases the data lines, and drops DataRdy

I/O device sees DataRdy go low and drops Ack

A Typical I/O System



Interfacing I/O Devices to the Processor, Memory, and OS

- The operating system acts as the interface between the I/O hardware and the program requesting I/O since

- Multiple programs using the processor share the I/O system
- I/O systems usually use interrupts which are handled by the OS
- Low-level control of an I/O device is complex and detailed

- Thus OS must handle interrupts generated by I/O devices and supply routines for low-level I/O device operations, provide equitable access to the shared I/O resources, protect those I/O devices/activities to which a user program doesn't have access and schedule I/O requests to enhance system throughput

OS must be able to give commands to the I/O devices

I/O device must be able to notify the OS about its status

Must be able to transfer data between the memory and the I/O device

Communication of I/O Devices and Processor

- How the processor directs the I/O devices

- Special I/O instructions

- Must specify both the device and the command

Memory-mapped I/O

- Portions of the high-order memory address space are assigned to each I/O device
- Read and writes to those memory addresses are interpreted as commands to the I/O devices

- Load/stores to the I/O address space can only be done by the OS

How I/O devices communicate with the processor

- Polling – the processor periodically checks the status of an I/O device (through the OS) to determine its need for service

- Processor is totally in control – but does all the work

- Can waste a lot of processor time due to speed differences

- Interrupt-driven I/O – the I/O device issues an interrupt to indicate that it needs attention